

# Parallel Programming

## Channels

Morgan Ericsson

# Today

- » Channels
- » Concurrency patterns

# Message passing

- » Processes communicate by sending messages
- » Synchronous or Asynchronous
- » Buffering
- » Kernel or Process

# Channels

```
1 func main() {  
2     strCh := make(chan string)  
3     go func() {  
4         strCh <- "Alice"  
5     }()  
6  
7     if res, ok := <-strCh; ok {  
8         fmt.Println("Hello", res)  
9     }  
10 }
```

# Channels

- » Channels send and receive messages
  - » `[channel] <- [value]` to send
  - » `[value], [status] <- [channel]` to receive
- » Channels are typed and can only send and receive messages of that type
- » A form of IPC

# CSP

- » Go's channels are based on Communicating Sequential Processes (CSP)
- » Idea that input and output need to be considered language primitives
- » Communication between processes
- » First-class support for modeling communication makes solving problems simpler and easier to comprehend

# CSP

- » `cardreader?cardimage`: read from cardreader and assign to cardimage
- » `lineprinter!lineimage`: send value of lineimage to lineprinter
- » `*[c:character; west?c → east!c]`: read all characters from west and pass them one by one to east. End when there are no more characters

# Go

- » Supports both shared-memory programming and primitives and CSP
- » CSP
  - » Channels models input and output
  - » Communication between goroutines



# Channels

# Go

```
1 func main() {  
2     strCh := make(chan string)  
3     go func() {  
4         strCh <- "Alice"  
5     }()  
6  
7     go func() {  
8         if res, ok := <-strCh; ok {  
9             fmt.Println("Hello", res)  
10        }  
11    }()  
12    time.Sleep(time.Second)  
13 }
```

# Channels

- » A language primitive
  - » A channel can be used as any other variable
- » Unbuffered or buffered
- » uni- or bidirectional
- » Composable

# Unbuffered

```
1 func main() {
2     tickCh := make(chan int)
3     go func() {
4         tmp := time.Now()
5         for i := 0; i < 5; i++ {
6             tickCh <- i
7             fmt.Println("Wrote after", time.Since(tmp))
8         }
9     }()
10
11     go func() {
12         for {
13             _ = <-tickCh
14             time.Sleep(time.Second)
15         }
16     }()
17     time.Sleep(8 * time.Second)
18 }
```

# Unbuffered

Wrote after 0s

Wrote after 1s

Wrote after 2s

Wrote after 3s

Wrote after 4s

# Unbuffered

- » Reads will block if the channel is empty
- » Writes will block if the channel is full
- » An unbuffered channel can hold a single value
- » Buffered channels can hold an arbitrary number of values
  - » Specified when created

# Buffered

```
1 func main() {
2     tickCh := make(chan int, 2)
3     go func() {
4         tmp := time.Now()
5         for i := 0; i < 5; i++ {
6             tickCh <- i
7             fmt.Println("Wrote after", time.Since(tmp))
8         }
9     }()
10
11     go func() {
12         for {
13             _ = <-tickCh
14             time.Sleep(time.Second)
15         }
16     }()
17     time.Sleep(8 * time.Second)
18 }
```

# Buffered

Wrote after 0s

Wrote after 0s

Wrote after 0s

Wrote after 1s

Wrote after 2s



# Bidirectional

```
1 func main() {  
2     mch := func() chan int {  
3         ch := make(chan int)  
4         go func() {  
5             _ = <-ch  
6         }()  
7         return ch  
8     }()  
9  
10    mch <- 1  
11 }
```

# Uni- and bidirectional

- » Channels are bidirectional by default (`chan`)
- » Can be implicitly converted to a unidirectional channel, e.g., by parameter or return type
- » Unidirectional channels are defined by adding `<-` before or after `chan`
  - » Position specifies direction

# Unidirectional

```
1 func main() {  
2     mch := func() <-chan int {  
3         ch := make(chan int)  
4         go func() {  
5             _ = <-ch  
6         }()  
7         return ch  
8     }()  
9  
10    mch <- 1  
11 }
```

./prog.go:14:2: invalid operation: cannot send to receive-only channel mch (variable of type <-chan int)

# Unidirectional

```
1 func main() {  
2     mch := func() chan<- int {  
3         ch := make(chan int)  
4         go func() {  
5             _ = <-ch  
6         }()  
7         return ch  
8     }()  
9  
10    mch <- 1  
11 }
```

# Closing channels

```
1 func main() {  
2     var wg sync.WaitGroup  
3     ch := make(chan int)  
4     wg.Add(1)  
5     go func() {  
6         defer wg.Done()  
7         st := time.Now()  
8         _ = <-ch  
9         fmt.Println("Channel closed after", time.Since(st))  
10    }()  
11    time.Sleep(5 * time.Second)  
12    close(ch)  
13    wg.Wait()  
14 }
```

# Closing channels

- » It is not possible to write to a closed channel
- » Reads will return already written values or nothing
- » The error return value indicates if the channel is open or closed

# Ranging over channels

```
1 gen := func () <-chan int {  
2     ch := make(chan int)  
3     go func() {  
4         defer close(ch)  
5         for i := 0; i<5; i++ {  
6             ch <- i  
7         }  
8     }()  
9     return ch  
10 }
```

# Ranging over channels

```
1 ch := gen()  
2 for r := range ch {  
3     fmt.Println("Got:", r)  
4 }
```



# Ranging over channels

- » `for ... range [channel]` is equivalent to reading in a loop and stopping when channel is drained and closed
- » Common pattern
  - » Just remember to close the channel

# Channels and state

## » Read

- » `nil` will **block**
- » *open* and *not empty* will **read** the value
- » *open* and *empty* will **block**
- » *closed* will **return** zero value and false
- » *send only* will not compile

# Channels and state

## » Write

- » `nil` will **block**
- » *open* and *full* will **block**
- » *open* and *not full* will **write** the value
- » *closed* will **panic**
- » *receive only* will not compile

# Channels and state

## » Close

- » `nil` will **panic**
- » *open* and *not empty* will **close** the channel but values are available until read
- » *open* and *empty* will **close** the channel. Any reads return the zero value
- » *closed* will **panic**
- » *receive only* will not compile

# Ownership

- » A channel is “owned” by the creator
  - » ownership can be passed
- » Owner should write to and close the channel
  - » again, can be passed
- » Use unidirectional channels to enforce
- » A goroutine does not own a read-only channel

# Why?

- » Writing to nil deadlocks
- » Closing a closed or nil channel panics
- » Writing to a closed channel panics
- » Owner's responsibility to manage
- » Consumer of a channel only needs to worry about
  - » reading from a closed channel
  - » blocking on reads

# Select

```
1 var c1, c2 <-chan int
2 select {
3     case <- c1:
4         // Do something
5     case <- c2:
6         // Do something else
7 }
```

# Select

- » Select is a switch for channel
- » Waits on multiple channels
  - » Unblocks and runs a case when possible



# Select

```
1 func main() {
2     st := time.Now()
3     ch := make(chan interface{})
4     go func() {
5         time.Sleep(5 * time.Second)
6         close(ch)
7     }()
8
9     fmt.Println("Blocking")
10    select {
11    case <-ch:
12        fmt.Println("Unblocked after", time.Since(st))
13    }
14 }
```

# Select

- » Closing a channel will unblock any operations waiting on it
- » Common way to signal goroutines that something has happened, e.g., that they should terminate

# Select

```
1 func main() {
2     c1 := make(chan interface{})
3     close(c1)
4     c2 := make(chan interface{})
5     close(c2)
6
7     var c1Count, c2Count int
8     for i := 1000; i >= 0; i-- {
9         select {
10            case <-c1:
11                c1Count++
12            case <-c2:
13                c2Count++
14        }
15    }
16
17    fmt.Println("c1Count:", c1Count, "c2Count:", c2Count)
18 }
```

# Select

- » Channels are not processed in order (unlike switch)
  - » would make it difficult to use select to manage multiple channels
- » The Go runtime randomly selects a channel if multiple are available
  - » Aims for a uniform distribution

# Runs

- » c1Count: 534 ; c2Count: 466
- » c1Count: 503 ; c2Count: 497
- » c1Count: 481 ; c2Count: 519
- » c1Count: 500 ; c2Count: 500
- » c1Count: 510 ; c2Count: 490

# Blocked

- » What if there are no available channels?
  - » selected blocked forever
- » Sometimes what we want, but not always
- » We can introduce a timeout or not block

# Unblocking after some time

```
1 func main() {
2     st := time.Now()
3     ch := make(chan interface{})
4     go func() {
5         time.Sleep(5 * time.Second)
6         close(ch)
7     }()
8
9     fmt.Println("Blocking")
10    select {
11    case <-ch:
12        fmt.Println("Unblocked after", time.Since(st))
13    case <-time.After(time.Second):
14        fmt.Println("Unblocked after", time.Since(st))
15    }
16 }
```

# Unblocking after some time

- » Note that `time.After` returns a channel that sends a value after some specified time
- » The default clause will be executed if there are no available values
  - » So, a select with default “never” blocks



# Unblocking

```
1 func main() {
2     st := time.Now()
3     ch := make(chan interface{})
4     go func() {
5         time.Sleep(5 * time.Second)
6         close(ch)
7     }()
8
9     fmt.Println("Blocking")
10    select {
11    case <-ch:
12        fmt.Println("Unblocked after", time.Since(st))
13    default:
14        fmt.Println("Unblocked after", time.Since(st))
15    }
16 }
```

# Patterns

# Patterns

- » Channels and goroutines can be used to design interesting solutions to concurrency problems
- » Can often avoid the need for shared state/memory
  - » State is communicated
  - » Synchronization from communication

# Done and for-select

```
1 func doSmtH(done <-chan interface{}) {  
2     for {  
3         select {  
4             case <-done:  
5                 return  
6             default:  
7             }  
8  
9             // Do some work  
10    }  
11 }
```

# Done and for-select

- » A done-channel can be used to tell goroutines to stop/exit
- » Default means we do not block and for that we check continuously
- » Note, `interface{}` (or `any`) specifies that the channel can accept any type
  - » But, does not matter since we never send anything to the channel

# Leaking Goroutines

```
1 func main() {  
2     doSmtH := func(str <-chan string) {  
3         for s := range str {  
4             fmt.Println(s)  
5         }  
6     }  
7     go doSmtH(nil)  
8     // Do other work ...  
9 }
```

# Leaking Goroutines

- » The for will never progress on a nil channel, so the function will never exit
  - » unless the program exits
- » A leaked goroutine
  - » similar to memory leaks
- » Can be a problem in long-running programs
  - » goroutines are cheap but not free
  - » resources not reclaimed until it exits

# Fixed (as long as we close done)

```
1 func doSmtH(str <-chan string, done <-chan interface{}) {  
2     for {  
3         select {  
4             case res, ok := <-str:  
5                 if !ok {  
6                     return  
7                 }  
8                 // do something with res  
9             case <-done:  
10                return  
11        }  
12    }  
13 }
```



# Done channels

- » When composing goroutines and channels we might end up with multiple done channels
- » Sufficient to wait for a single done message, so reduces to an or
- » So, we need or over multiple channels

# Or channel

```
1 func or(ch1, ch2 <-chan interface{}) <-chan interface{} {
2     orDone := make(chan interface{})
3     go func() {
4         defer close(orDone)
5         select {
6             case <-ch1:
7             case <-ch2:
8         }
9     }()
10    return orDone
11 }
```

# Or channel

- » The or function takes two channels and returns a channel
- » The returned channel will close when either of the two channels unblocks
- » Can use this idea and varadic functions to create an or that waits on an arbitrary number of done channels
- » Note that this pattern cannot be used to multiplex channels

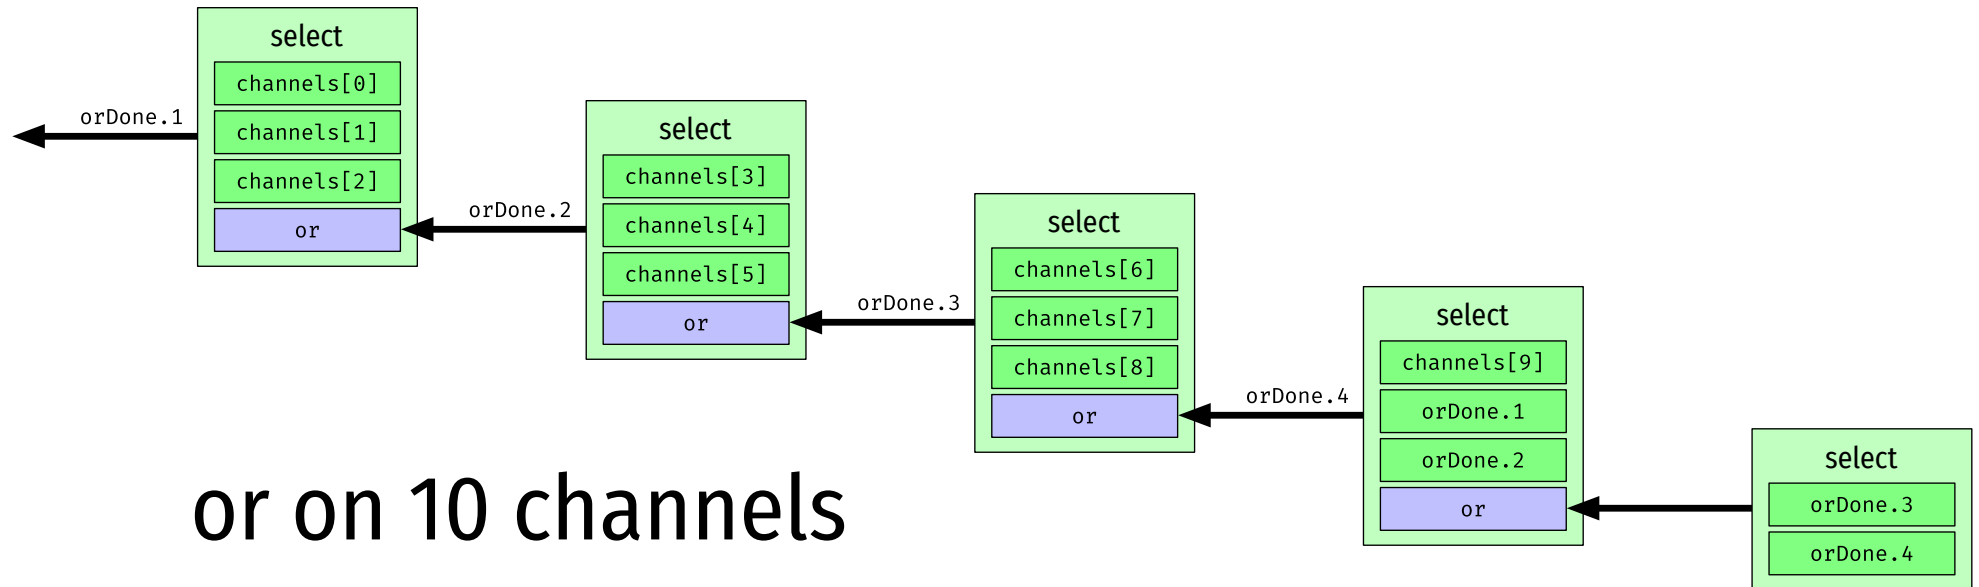
# Arbitrary number of channels?

```
1 func or(channels ...<-chan interface{}) <-chan interface{} {  
2     switch len(channels) {  
3     case 0: return nil  
4     case 1: return channels[0]  
5     }  
6  
7     //...
```

# Arbitrary number of channels?

```
1  // ...
2  orDone := make(chan interface{})
3  go func() {
4      defer close(orDone)
5      switch len(channels) {
6      case 2:
7          select {
8              case <-channels[0]:
9              case <-channels[1]:
10          }
11      default:
12          select {
13              case <-channels[0]:
14              case <-channels[1]:
15              case <-channels[2]:
16              case <-or(append(channels[3:], orDone)...):
17          }
18      }
19  }()
20  return orDone
21 }
```

# Arbitrary number of channels?



# Example

```
1 sig := func(after time.Duration) <-chan interface{} {
2     ch := make(chan interface{})
3     go func() {
4         defer close(ch)
5         time.Sleep(after)
6     }()
7     return ch
8 }
9
10 start := time.Now()
11 <-or(
12     sig(2 * time.Hour),
13     sig(5 * time.Minute),
14     sig(time.Second),
15     sig(time.Hour),
16     sig(time.Minute),
17 )
18 fmt.Println("Done after", time.Since(start))
```

# isPrime (from Wikipedia)

```
1 func isPrime(num int) bool {  
2     if num > 1 && num <= 3 {  
3         return true  
4     }  
5     if num <= 1 || num%2 == 0 || num%3 == 0 {  
6         return false  
7     }  
8  
9     for i := 5; i*i <= num; i += 6 {  
10        if num%i == 0 || num%(i+2) == 0 {  
11            return false  
12        }  
13    }  
14    return true  
15 }
```



# Find primes

```
1 func main() {  
2     for i:=0;i<1000000;i++ {  
3         if isPrime(i) {  
4             fmt.Println("Got", i)  
5         }  
6         i += 1  
7     }  
8 }
```

# Fork-join

```
1 ch := make(chan int)
2 for i := 0; i < 10; i++ {
3     go func(lo, hi int) {
4         for i := lo; i < hi; i++ {
5             if isPrime(i) {
6                 ch <- i
7             }
8         }
9     }(i*100000, (i+1)*100000)
10 }
11
12 for p := range ch {
13     fmt.Println(p)
14 }
```

Does not work: “fatal error: all goroutines are asleep - deadlock!”

# Why?

- » All the goroutines will exit when they have checked their range
- » The channel is still open, so the for will expect more values
- » But there will be no more values, so progress is not possible!
- » How can we close the channel?

# Fork-join

```
1 var wg sync.WaitGroup
2 ch := make(chan int)
3 for i := 0; i < 10; i++ {
4     wg.Add(1)
5     go func(lo, hi int) {
6         defer wg.Done()
7         for i := lo; i < hi; i++ {
8             if isPrime(i) {
9                 ch <- i
10            }
11        }
12    }(i*100000, (i+1)*100000)
13 }
14
15 for p := range ch {
16     fmt.Println(p)
17 }
```

# Fork-join

- » Channel is basically a thread-safe queue
- » We split based on benchmarks (or gut feeling)
- » Can we try a different approach?

# Pipelining

```
1 func filterPrime(ins <-chan int) <-chan int {
2     outCh := make(chan int)
3     go func() {
4         for i := range ins {
5             if isPrime(i) {
6                 outCh <- i
7             }
8         }
9     }()
10
11     return outCh
12 }
```

# Pipelining

```
1 func genInts(done <-chan interface{}) <-chan int {
2     outCh := make(chan int)
3     go func() {
4         i := 1
5         for {
6             select {
7                 case <-done:
8                     close(outCh)
9                     return
10                case outCh <- i:
11                }
12                i += 1
13            }
14        }()
15     return outCh
16 }
```

# Remember


- » Note that both functions own their channels
  - » Returns a read-only version
- » Minor issue in filterPrimes, can leak if channel is never closed
  - » Should fix with the done-channel pattern



# Pipelining

```
1 func main() {  
2     var doneCh chan interface{}  
3  
4     go func() {  
5         time.Sleep(5 * time.Second)  
6         close(doneCh)  
7     }()  
8  
9     intS := genInts(doneCh)  
10    divCh := filterPrime(intS)  
11    for r := range divCh {  
12        fmt.Println("Got", r)  
13    }  
14    fmt.Println("Done")  
15 }
```

# Pipelining



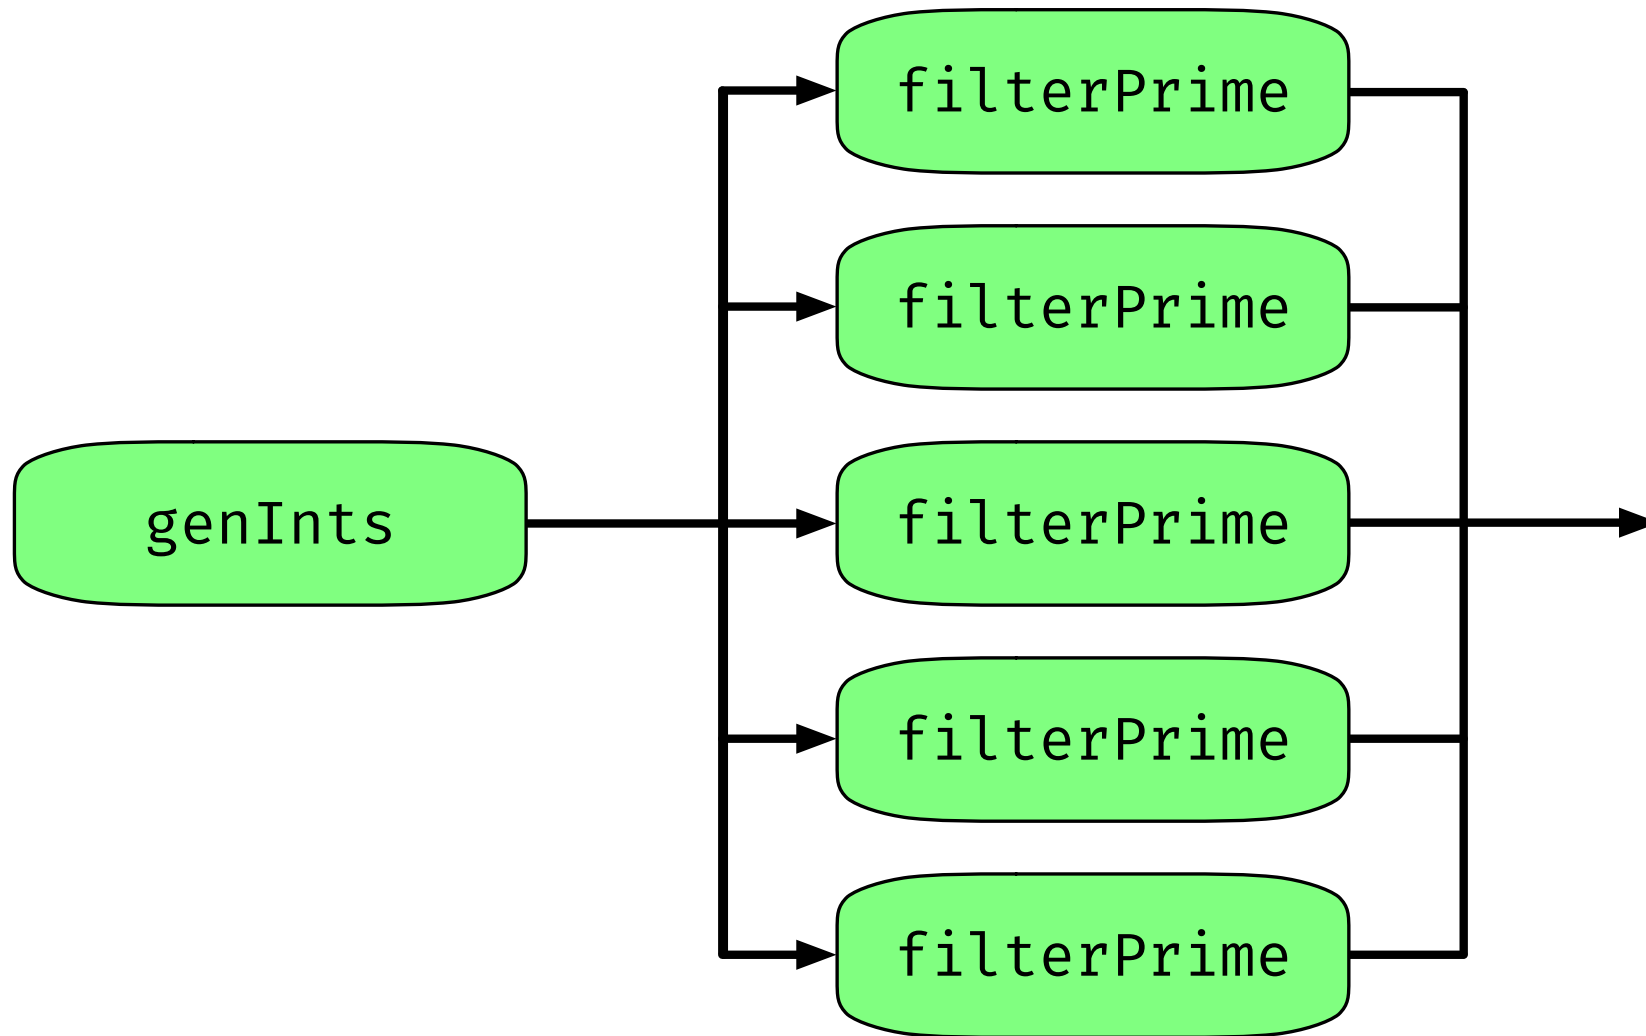
# Pipelining

- » Our solution drains the int channel
- » We could pass the done channel to `filterPrime` and stop without necessarily draining
- » Common pattern in pipelines that might not drain the channel
- » We risk leaking the goroutine!

# Pipelining

- » In this case, no meaningful concurrency
  - » Generation and checking split
- » However, quite easy to add concurrency
  - » By “fanning out” the pipeline

# Fan-out



# Fan-out

- » If shared output channel, when and where do we close it?
- » If N output channels, how do we handle these?
  - » Drain each in a loop?
- » Fan-in
  - » Similar to or, but forwards values

# Fan-in


```
1 func multiplexer(channels ...<-chan int) <-chan int {
2     var wg sync.WaitGroup
3     mpOut := make(chan int)
4
5     for _, ch := range channels {
6         wg.Add(1)
7         go func(c <-chan int) {
8             defer wg.Done()
9             for i := range c {
10                 mpOut <- i
11             }
12         }(ch)
13     }
14
15     go func() {
16         wg.Wait()
17         close(mpOut)
18     }()
19
20     return mpOut
21 }
```

# Primes

```
1 func main() {
2     var doneCh chan interface{}
3     chs := make([]<-chan int, 5)
4
5     go func() {
6         time.Sleep(5 * time.Second)
7         close(doneCh)
8     }()
9
10    ints := genInts(doneCh)
11    for i:=0;i<5;i++ {
12        chs[i] = filterPrime(ints)
13    }
14    pCh := multiplexer(chs...)
15    for r := range pCh {
16        fmt.Println("Got", r)
17    }
18    fmt.Println("Done")
19 }
```



# Primes



# Fixing the potential leak

```
1 func filterPrime(ins <-chan int, done <-chan interface{}) <-chan int {
2     outCh := make(chan int)
3     go func() {
4         defer close(outCh)
5         for {
6             select {
7                 case <-done:
8                     return
9                 case i, ok := <-ins:
10                    if !ok { return }
11            }
12            if isPrime(i) {
13                outCh <- i
14            }
15        }
16    }()
17
18    return outCh
19 }
```

# Annoying!

- » We cannot range over the channel
- » Easy to make mistakes
- » Multiplex the two channels
  - » done closes the channel we read from
- » Allows a “simple” for range

# Or-done

```
1 func orDone(done, ch <-chan interface{}) <-chan interface{} {
2     outCh := make(chan interface{})
3     go func() {
4         defer close(outCh)
5         for {
6             select {
7                 case <-done:
8                     return
9                 case v, ok := <-ch:
10                    if !ok { return }
11                    select {
12                        case outCh <- v:
13                        case <-done:
14                        }
15                    }
16                }
17            }()
18            return outCh
19 }
```

# New version

```
1 func filterPrime(ins <-chan int, done <-chan interface{}) <-chan int {
2     outCh := make(chan int)
3     go func() {
4         for i := range orDone(ins, done) {
5             if isPrime(i) {
6                 outCh <- i
7             }
8         }
9     }()
10
11     return outCh
12 }
```

# Channels

- » We can do a lot more with channels
  - » A tee-channel that splits for, e.g., logging
  - » ...
- » We can pass channels over channels
  - » The multiplexer could be rewritten to take a channel that contains channels
  - » And bridge these

# Concurrency in Go

# Remember

Concurrency is a property of the code; parallelism is a property of the running program.



# Concurrency in Go

Do not communicate by sharing memory. Instead, share memory by communicating.

# Always?

- » Channels should be your first choice
- » Unless you need to protect some internal state
- » If channels are not performant enough, consider shared memory
  - » But try to refactor first

# Other languages

- » Similar ideas work
- » Tasks (however they are formulated) and thread-safe queues

# Next time

» Deadlocks

