

Parallel Programming

Parallel algorithms 1

Morgan Ericsson

Today

- » Prefix sum/scan
- » Sorting

Sum and scan

Remember: concurrent vs parallel

- » Concurrency is about structure
- » Parallelism is about execution
- » For the algorithms so be beneficial
 - » Parallelism to execute
 - » Concurrency to structure

Computing the sum

```
1 var s int
2 l := []int{3, 5, 2, 5, 7, 9, 4, 6}
3
4 for _, val := range l {
5     s += val
6 }
7
8 fmt.Println("sum:", s)
```

Computing the sum

- » Remember the algorithms course
- » 8 add operations
- » $O(N)$
- » Can we do it in fewer with parallelism?


Aside: Not just the sum

- » The previous example is a pattern
- » We can use any binary commutative operation
 - » sum, product, max, ...

Aside: Not just the sum

```
1 s := 1
2 m := math.MinInt
3 l := []int{3, 5, 2, 5, 7, 9, 4, 6}
4
5 for _, val := range l {
6     s *= val
7     m = max(m, val)
8 }
9
10 fmt.Printf("prod: %d\nmax: %d", s, m)
```


How many operations do we need?



Aside: PRAM

- » Parallel Random Access Machine
- » Shared-memory multiprocessor model
- » Unlimited number of processor with unlimited local and shared memory
- » Each processor knows its ID
- » Inputs and outputs are placed in shared memory
- » Each instruction takes unit time
- » Instructions are synchronised across processors

PRAM

- » Unfeasible model
 - » The interconnection network between processors and memory would require a very large area
 - » The message routing would require time proportional to the network size
- » Theoretical model
- » Algorithm designers can forget about communication and focus on the parallel computation

Implementation

```
1 l := []int{3, 5, 2, 5, 7, 9, 4, 6}
2 tmp := int(math.Log2(float64(len(l))))
3
4 for j := 1; j < tmp+1; j++ {
5     for k := 0; k < len(l); k++ {
6         if (k+1)%(1<<j) == 0 {
7             l[k] = l[k-1<<(j-1)] + l[k]
8         }
9     }
10 }
```


l: [3 8 2 15 7 16 4 41]

sum: 41

<< and >>

- » Shifts the bits in a number left or right a number of times
 - » `a := 1 ; a <<= 3`
 - » `0001` shifted three times to the left becomes `1000`
- » Can be used to implement multiplication and division by two on integers
- » Used here since explicit type conversion quickly becomes unreadable

Implementation



Implementation (makes little sense)

```
1  var wg sync.WaitGroup
2  l := []int{3, 5, 2, 5, 7, 9, 4, 6}
3  tmp := int(math.Log2(float64(len(l))))
4
5  for j := 1; j < tmp+1; j++ {
6      for k := 0; k < len(l); k++ {
7          wg.Add(1)
8          go func(x, y int) {
9              defer wg.Done()
10             if (x+1)%(1<<y) == 0 {
11                 l[x] = l[x-1<<(y-1)] + l[x]
12             }
13         }(k, j)
14     }
15 }
16 wg.Wait()
```

Many issues

- » We assume one processor per element
 - » So 8 processors for 8 elements
- » We cover races/sync by starting and stopping threads
 - » This is equivalent to a barrier
 - » but slower...

A more practical approach

1. Split the array into one part per thread (fork)
2. Compute the sum for each part
3. Sum the sums of the parts (join)

A more practical approach

```
1 func summation(lst []int, wid int, numworks int, out chan<- int) {
2     beg := len(lst) / numworks * wid
3     end := len(lst) / numworks * (wid + 1)
4
5     if wid == numworks - 1 {
6         end = len(lst)
7     }
8
9     var ls int
10    for i := beg; i < end; i++ {
11        ls += lst[i]
12    }
13
14    out <- ls
15 }
```

A more practical approach

```
1 func main() {
2     ch := make(chan int)
3     l := make([]int, 10000)
4     var tmp int
5     for i := 0; i < 10000; i++ {
6         l[i] = i
7     }
8
9     for i := 0; i < 8; i++ {
10        go summation(l, i, 8, ch)
11    }
12
13    var gs int
14    for i := 0; i < 8; i++ {
15        gs += <-ch
16    }
17 }
```

So much setup in summation!

```
1 func summation2(lst []int, wid int, numworks int, out chan<- int) {  
2     var ls int  
3     for i := wid; i < len(lst); i += numworks {  
4         ls += lst[i]  
5     }  
6  
7     out <- ls  
8 }
```

Difference?

3	5	2	5	7	9	4	6
---	---	---	---	---	---	---	---

8	7	16	10
---	---	----	----

3	5	2	5	7	9	4	6
---	---	---	---	---	---	---	---

10	14	6	11
----	----	---	----

Prefix scan

- » Computes all partial sums of a vector of values
- » Inclusive and exclusive
 - » Exclusive does not include “current” value
- » Just as parallel sum, any associative combining operation can be used

Prefix scan

Original

3	5	2	5	7	9	4	6
---	---	---	---	---	---	---	---

Inclusive

3	8	10	15	22	31	35	41
---	---	----	----	----	----	----	----


Exclusive

0	3	8	10	15	22	31	35
---	---	---	----	----	----	----	----

Again, simple serial


```
1 lst := []int{3, 5, 2, 5, 7, 9, 4, 6}
2 ps := make([]int, len(lst))
3
4 ps[0] = lst[0]
5 for i := 1; i < len(lst); i++ {
6     ps[i] = ps[i-1] + lst[i]
7 }
```


Unrealistic parallel version



```
1 ll := int(math.log2(float64(len(l))))
2 for j:=1;j<ll+1;j++ {
3     // This should run on
4     // k processors
5     for k:=0;k<len(l);k++ {
6         go func(k int) {
7             if k >= (1<<j - 1) {
8                 l[k] = l[k] + \
9                     l[k - (1<<j-1)]
10            }
11        }(k)
12    }
13 }
```

A practical approach



Prefix scan

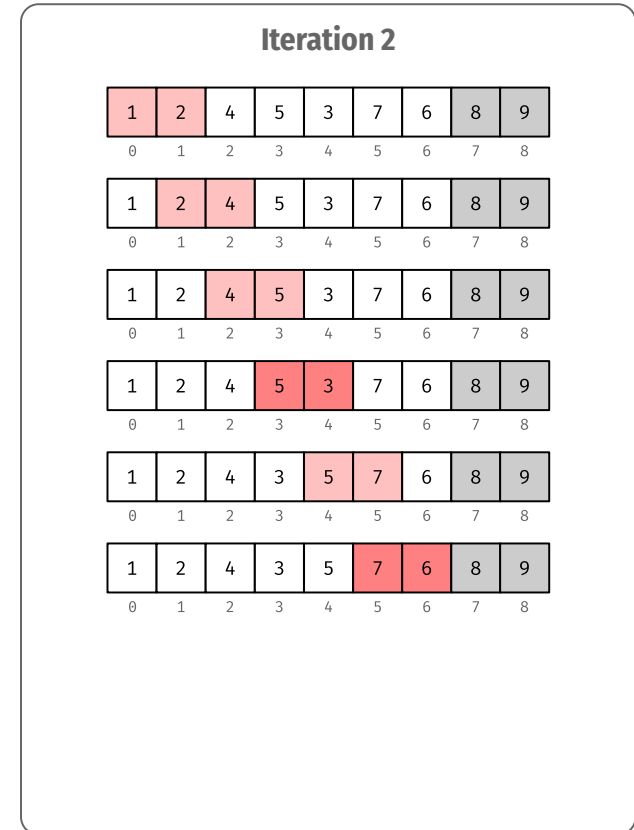
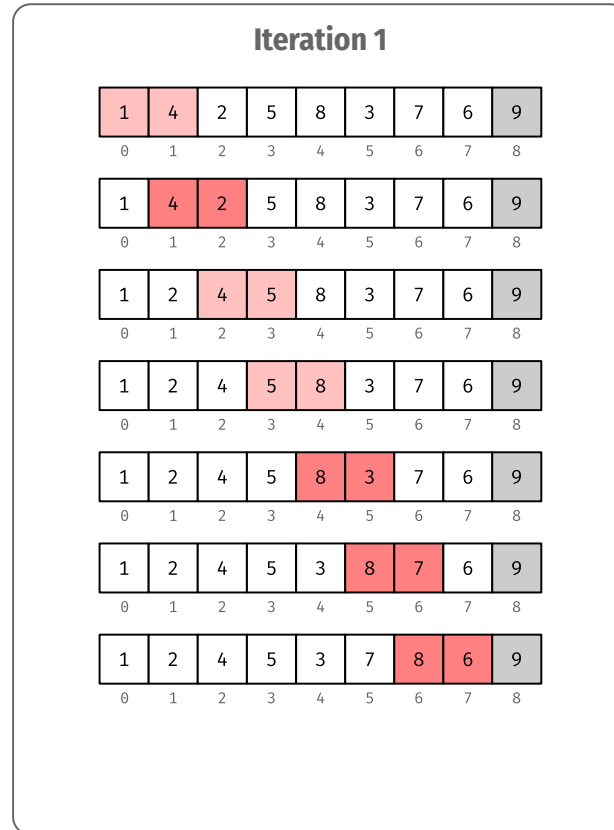
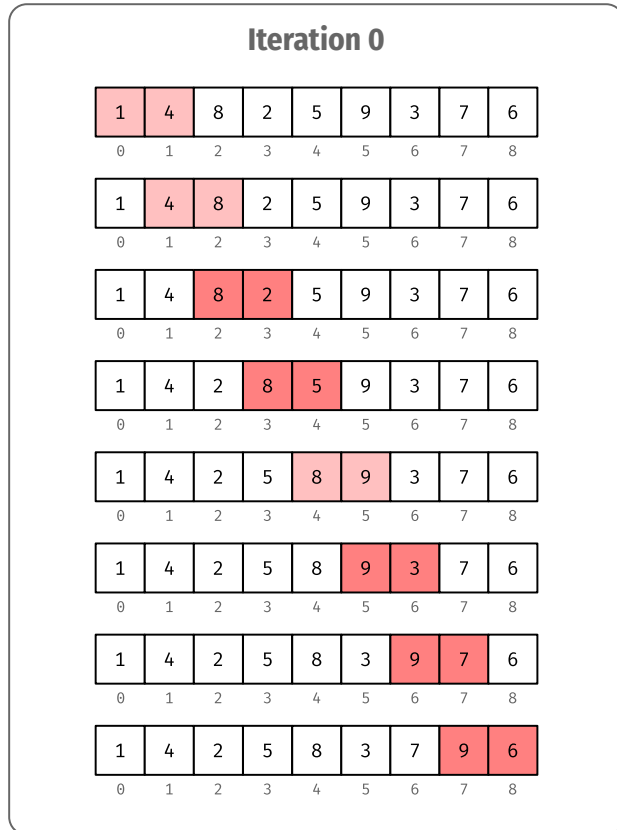
```
1 func pscan(lst *[]int, wid int, ws int, sz int, ch chan int) {
2     beg := sz / ws * wid
3     end := sz / ws * (wid + 1)
4     if wid == ws-1 {
5         end = sz
6     }
7
8     for i := beg + 1; i < end; i++ {
9         (*lst)[i] = (*lst)[i-1] + (*lst)[i]
10    }
11
12    ch <- (*lst)[end-1]
13    tmp := <-ch
14
15    for i := beg; i < end; i++ {
16        (*lst)[i] = (*lst)[i] + tmp
17    }
18 }
```

Main

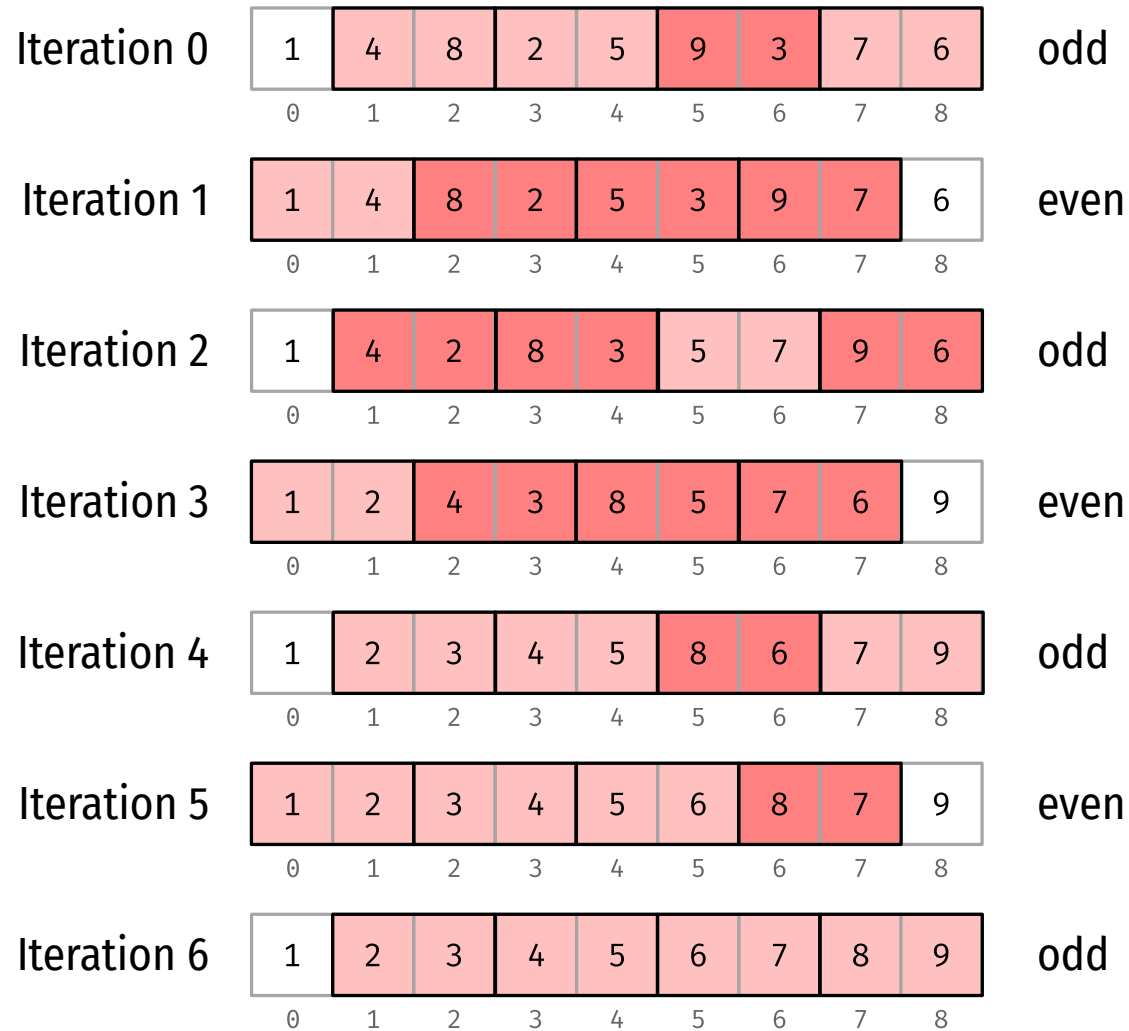
```
1 l := []int{3, 5, 2, 6, 7, 9, -4, 6, 7, -3, 1, 7, 6, 8, -1, 2}
2 chs := make([]chan int, 4)
3 var wg sync.WaitGroup
4
5 for i := 0; i < 4; i++ {
6     chs[i] = make(chan int)
7     wg.Add(1)
8     go func(ch chan int) {
9         defer wg.Done()
10        pscan(&l, i, 4, len(l), ch)
11    }(chs[i])
12 }
13
14 var rc int
15 for i := 0; i < 4; i++ {
16     tmp := <-chs[i]
17     chs[i] <- rc
18     rc += tmp
19 }
20
21 wg.Wait()
```

Sorting

Remember Bubble sort?



We can do this in parallel!



Odd-even transposition sort (serial)

```
1 l := []int{3, 5, 2, 6, 7, 9, -4, 6, 7, -3, 1, 7, 6, 8, -1, 2}
2 exch := true
3 var offs int
4
5 for exch {
6     exch = false
7     for i := 1 + offs; i < len(l); i += 2 {
8         if l[i-1] > l[i] {
9             l[i-1], l[i] = l[i], l[i-1]
10            exch = true
11        }
12    }
13    offs = (offs + 1) % 2
14 }
```


Parallel version

- » Split the array into N chunks (where N is the number of threads)
- » Sorting on each chunk in parallel
- » Challenges
 - » Did we swap? Shared
 - » The two phases (odd/even) should not overlap

Parallel version

```
1 func oets(lst *[]int, wid int, ws int, offs *int, exch *bool) {
2     beg := len(*lst) / ws * wid
3     end := (len(*lst) / ws * (wid + 1)) + 1
4     if wid == ws-1 {
5         end = len(*lst)
6     }
7     for exch {
8         exch = false
9         for i := beg + 1 + offs; i < end; i += 2 {
10             if (*lst)[i-1] > (*lst)[i] {
11                 (*lst)[i-1], (*lst)[i] = (*lst)[i], (*lst)[i-1]
12                 exch = true
13             }
14         }
15         ch <- exch
16         exch = <-ch
17         offs = (offs + 1) % 2
18     }
19 }
```

Fails

- » Problem with shared variables? Yes
 - » Assignment in Python is atomic
 - » All write the same value, so no race with flag
 - » But offset is problematic!
- » The threads can (and will) get out of sync
 - » Some threads might not even enter the while-loop

Second parallel version

```
1 func oets(lst []*int, wid int, ws int, ch chan bool) {
2     exch := true
3     var offs int
4     beg := len(*lst) / ws * wid
5     end := len(*lst) / ws * (wid + 1)
6     if wid == ws-1 {
7         end = len(*lst)
8     }
9     for exch {
10        exch = false
11        for i := beg + 1 + offs; i < end; i += 2 {
12            if (*lst)[i-1] > (*lst)[i] {
13                (*lst)[i-1], (*lst)[i] = (*lst)[i], (*lst)[i-1]
14                exch = true
15            }
16        }
17        ch <- exch
18        exch = <-ch
19        offs = (offs + 1) % 2
20    }
21 }
```

Second parallel version

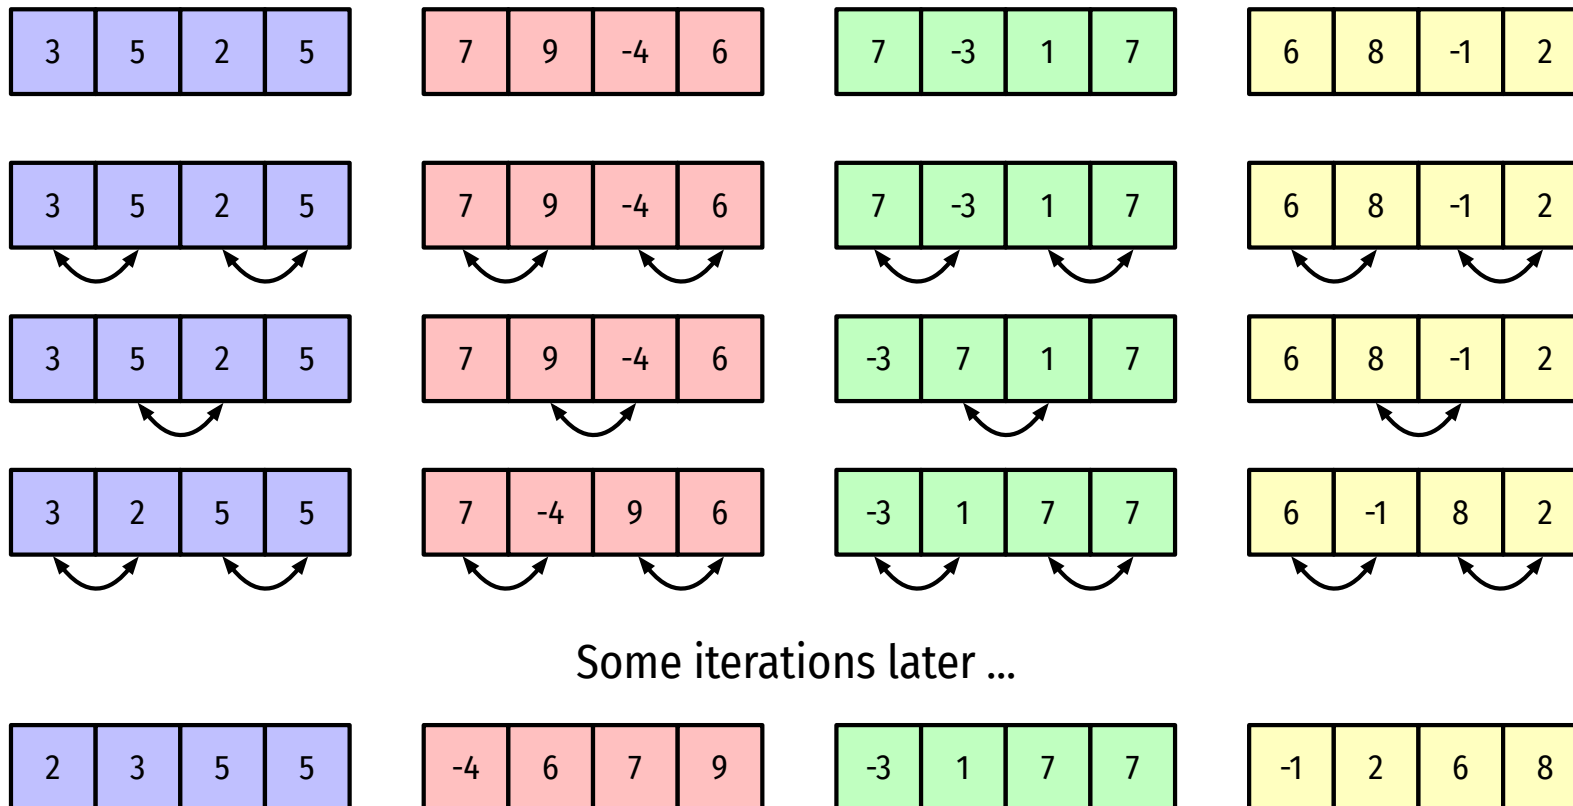
```
1 l := []int{3, 5, 2, 5, 7, 9, -4, 6, 7, -3, 1, 7, 6, 8, -1, 2}
2 chs := make([]chan bool, 4)
3
4 for i := 0; i < 4; i++ {
5     chs[i] = make(chan bool)
6     go func(ch chan bool) {
7         oets(&l, i, 4, ch)
8     }(chs[i])
9 }
10
11 // ...
```

Second parallel version

```
1 // ...
2 for {
3     var gexch bool
4     for i := 0; i < 4; i++ {
5         if tmp := <-chs[i]; tmp {
6             gexch = true
7         }
8     }
9     if gexch {
10         for i := 0; i < 4; i++ {
11             chs[i] <- gexch
12         }
13     } else {
14         break
15     }
16 }
```




Still fails

3	5	2	5	7	9	-4	6	7	-3	1	7	6	8	-1	2
---	---	---	---	---	---	----	---	---	----	---	---	---	---	----	---



Solution

3	5	2	5	7	9	-4	6	7	-3	1	7	6	8	-1	2
---	---	---	---	---	---	----	---	---	----	---	---	---	---	----	---



Solution

```
1 func oets(lst *[]int, wid int, ws int, ch chan bool) {
2     exch := true
3     var offs int
4     beg := len(*lst) / ws * wid
5     end := (len(*lst) / ws * (wid + 1) + )
6     if wid == ws-1 {
7         end = len(*lst)
8     }
9     for exch {
10        exch = false
11        for i := beg + 1 + offs; i < end; i += 2 {
12            if (*lst)[i-1] > (*lst)[i] {
13                (*lst)[i-1], (*lst)[i] = (*lst)[i], (*lst)[i-1]
14                exch = true
15            }
16        }
17        ch <- exch
18        exch = <-ch
19        offs = (offs + 1) % 2
20    }
21 }
```

Remember Quicksort

```
1 def _sort(a:list[int], lo:int, hi:int) -> None:
2     if hi <= lo:
3         return
4     j = _partition(a, lo, hi)
5     _sort(a, lo, j - 1)
6     _sort(a, j + 1, hi)
7
8 def sort(a:list[int]) -> None:
9     _sort(a, 0, len(a) - 1)
```

Remember Quicksort

- » We will do this one in Python
- » To see how things work in Python
- » And to leave room for you to implement it in Go!

Should be easy to parallelize?

```
1 def _sort(a:list[int], lo:int, hi:int) -> None:
2     if hi <= lo:
3         return
4     j = self._partition(a, lo, hi)
5     t1 = Thread(target=_sort, args=(a, lo, j - 1))
6     t2 = Thread(target=_sort, args=(a, j + 1, hi))
7     t1.start()
8     t2.start()
9     t1.join()
10    t2.join()
```

Terrible approach

- » How many threads will be created? Many!
- » Why is that a problem?
 - » Limit the number of threads
 - » High overhead from creation
- » Generally more difficult to deal with recursive algorithms

Make quicksort iterative

```
1  from queue import Queue
2
3  class IQS:
4      def __init__(self):
5          self.Q = Queue()
6
7      def _sort(self):
8          while not self.Q.empty():
9              lo, hi = self.Q.get()
10             if lo < hi:
11                 mid = self._partition(lo, hi)  # Unchanged
12
13                 self.Q.put((lo, mid - 1))
14                 self.Q.put((mid + 1, hi))
```

Make quicksort iterative

```
1 def sort(self:IQS, lst):  
2     self.lst = lst  
3     self.Q.put((0, len(self.lst) - 1))  
4     self._sort()
```

Easier to parallelize?

- » We use a queue instead of recursion
- » Each call to partition is independent, so ...
- » ... we can run them in parallel
- » Rather than spawning threads for the calls, ...
- » ... existing threads grab “jobs” from the queue
- » Remember, channels in Go are essentially concurrent queues!

New `_sort`

```
1 def _sort(self):
2     while True:
3         self.qsem.acquire():
4         lo, hi = self.Q.get()
5         if lo < hi:
6             mid = self._partition(lo, hi)
7             self.Q.put((lo, mid - 1))
8             self.Q.put((mid + 1, hi))
9             self.qsem.release(2)
```

New sort

```
1 def sort(self, l):
2     self.lst = l
3     self.Q.put((0, len(self.lst) - 1))
4     self.qsem.release()
5     ts = [Thread(target=self._sort)
6           for _ in range(self.nt)]
7     for t in ts:
8         t.start()
9     for t in ts:
10        t.join()
```

New `init`

```
1 def __init__(self):  
2     self.Q = Queue()  
3     self.qsem = Semaphore(0)
```

Unfortunately ...

- » ... it is a bit more complicated
- » How do we know when we are done?
- » We should not make timing assumptions
 - » e.g., end when there is nothing in the queue for a while
- » So, we need a flag ...
- » ... that is set once all elements are sorted

Adding a count

```
1  def _sort(self):
2      while not self.done:
3          self.qsem.acquire():
4          lo, hi = self.Q.get()
5          if lo < hi:
6              mid = self._partition(lo, hi)
7              self.sortcnt += 1
8              self.Q.put((lo, mid - 1))
9              self.Q.put((mid + 1, hi))
10             self.qsem.release(2)
11         if lo == hi:
12             self.sortcnt += 1
13             if self.sortcnt == self.N:
14                 self.done = True
```

Problem 1

```
1  def _sort(self):
2      while not self.done:
3          self.qsem.acquire():
4          lo, hi = self.Q.get()
5          if lo < hi:
6              mid = self._partition(lo, hi)
7              self.sortcnt += 1
8              self.Q.put((lo, mid - 1))
9              self.Q.put((mid + 1, hi))
10             self.qsem.release(2)
11         if lo == hi:
12             self.sortcnt += 1
13             if self.sortcnt == self.N:
14                 self.done = True
```

Problem 1

```
1 self.cntlock.acquire()  
2 self.sortcnt += 1  
3 self.cntlock.release()
```

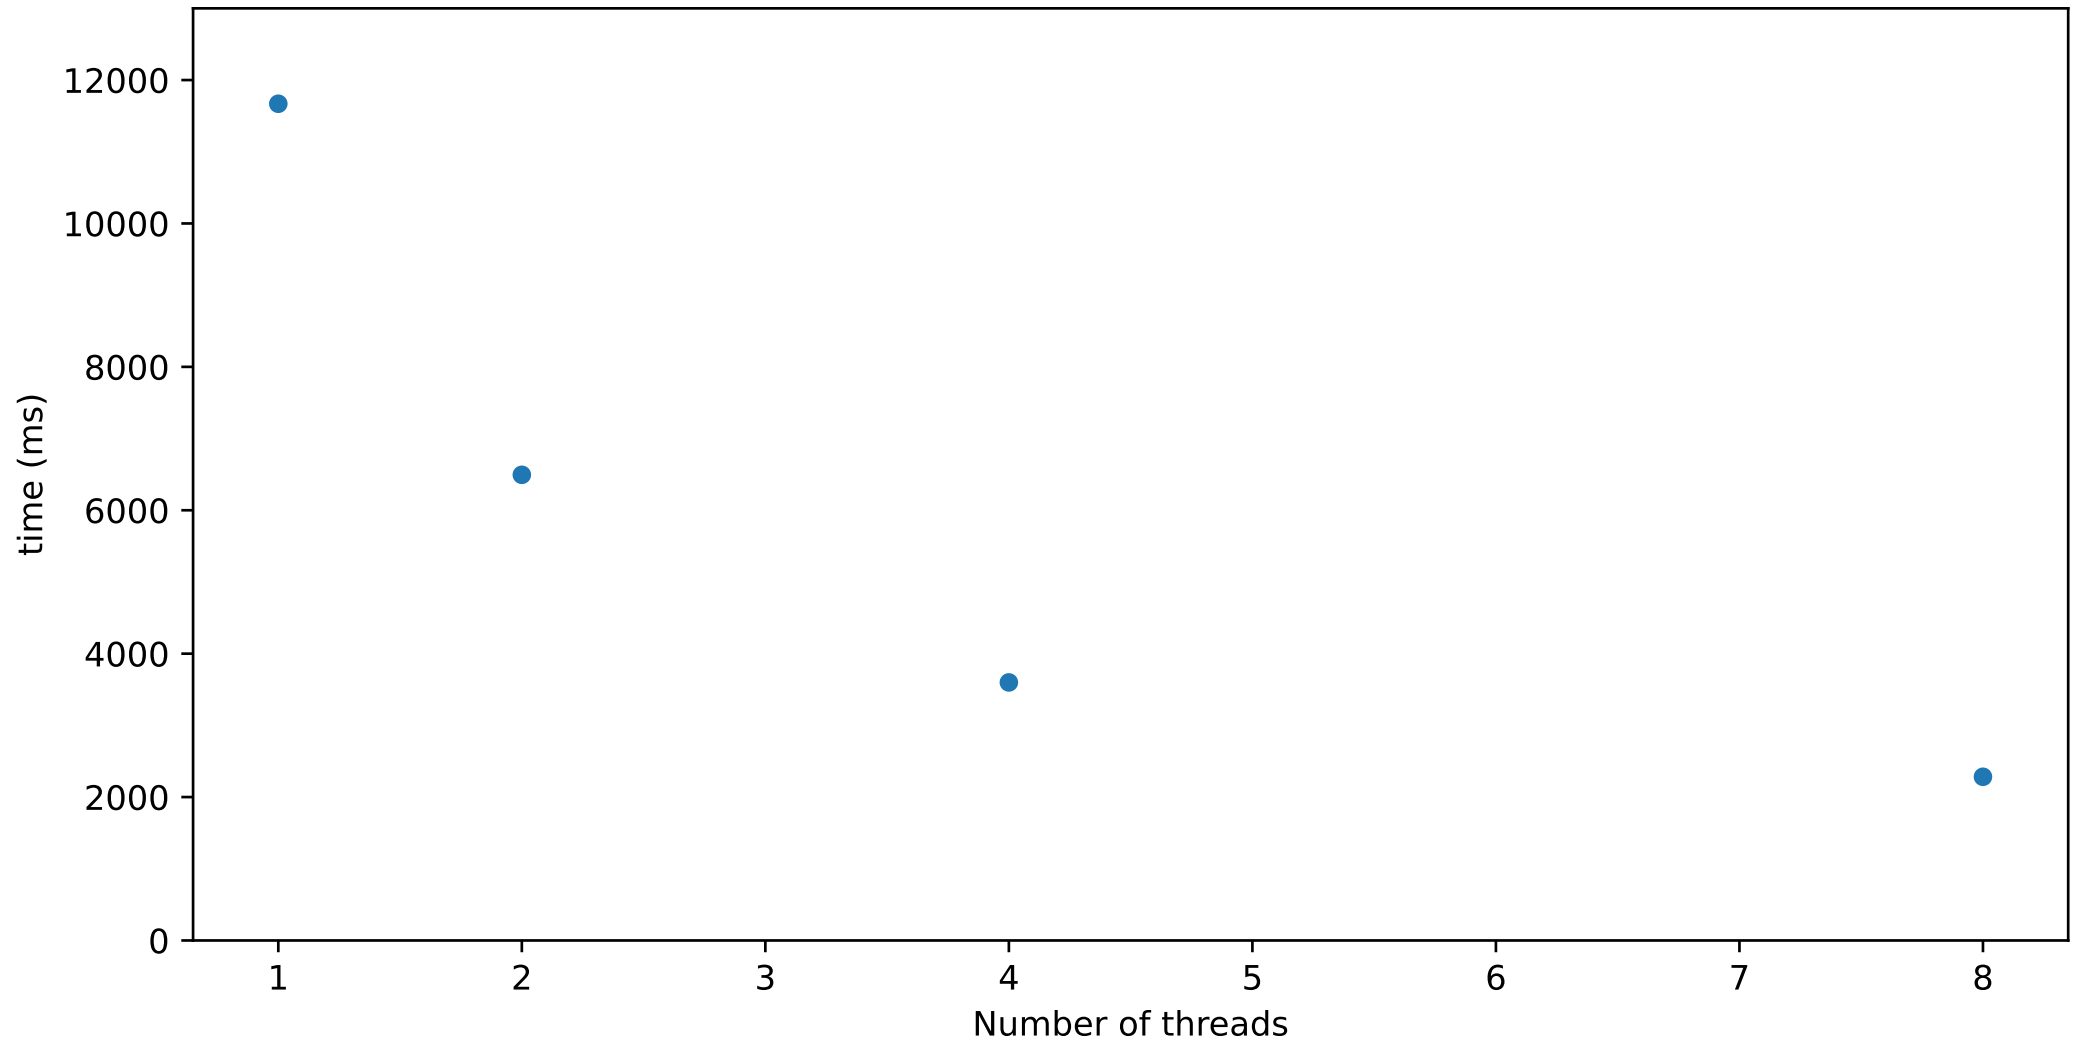
Problem 2

```
1  def _sort(self):
2      while not self.done:
3          self.qsem.acquire():
4          lo, hi = self.Q.get()
5          if lo < hi:
6              mid = self._partition(lo, hi)
7              self.sortcnt += 1
8              self.Q.put((lo, mid - 1))
9              self.Q.put((mid + 1, hi))
10             self.qsem.release(2)
11         if lo == hi:
12             self.sortcnt += 1
13             if self.sortcnt == self.N:
14                 self.done = True
```


Problem 2

```
1 while True:
2     while not self.qsem.acquire(timeout=10):
3         if self.done:
4             return
```

Scaling (Java version)



What did we do?

- » N threads and a queue
- » Basically an executor service

So, rewrite?

```
1 class CCQsort:
2     def __init__(self, nt):
3         self.exec = ThreadPoolExecutor(nt)
4
5     def sort(self, l):
6         self.lst = l
7         r = self.exec.submit(self._sort, 0, len(l) - 1)
8
9     def _sort(self, lo, hi):
10        if lo < hi:
11            mid = self._partition(lo, hi)
12            r1 = self.exec.submit(self._sort, lo, mid - 1)
13            r2 = self.exec.submit(self._sort, mid + 1, hi)
```

Does not work

- » Submit is not blocking, so when `sort(lst)` returns, the list is not sorted
- » We need to get sort to block ...
- » Or get something we can wait for ...

We had the same problem before!

```
1 def _sort(self, lo, hi):
2     if lo < hi:
3         mid = self._partition(lo, hi)
4         self.tl.acquire()
5         self.tc += 1
6         self.tl.release()
7         self.exec.submit(self._sort, lo, mid - 1)
8         self.exec.submit(self._sort, mid + 1, hi)
9     elif lo == hi:
10        self.tl.acquire()
11        self.tc += 1
12        tmp = self.tc
13        self.tl.release()
14        if tmp == self.N:
15            self.tsig.set()
```

We had the same problem before!

```
1 def sort(self, l):  
2     self.lst = l  
3     self.N = len(l)  
4     self.exec.submit(self._sort, 0, len(l) - 1)  
5     self.tsig.wait()
```


Signal/Event

- » We can replace the signal/Event with a lock
- » Acquire/lock when created
- » Try to lock in main thread (which will block)
- » Unlock when done (so main thread can lock)

What about partition?

- » Can we partition concurrently?
- » Single pivot, need to move “all” elements
- » So, need more than split into N parts and process


What about partition?



What about partition? (serial)

```
1 def ppartition(A, x):
2     B, lt, gt = [0] * len(A), [0] * len(A), [0] * len(A)
3
4     for i in range(len(A)):
5         B[i] = lst[i]
6         if B[i] < x: lt[i] = 1
7         if B[i] > x: gt[i] = 1
8
9     lt = prefixsum(lt)
10    gt = prefixsum(gt)
11    k = lt[n - 1]
12    lst[k] = x
13
14    for i in range(n):
15        if B[i] < x:
16            lst[lt[i] - 1] = B[i]
17        elif B[i] > x:
18            lst[k + gt[i]] = B[i]
19
20    return k
```

What about partition?



What about partition?

- » We can make the two for loops parallel
 - » Chunk or different offsets
- » We can use the parallel prefix scan
- » Need a few synchronization points with barriers

Summary (part 1)

New designs/tools

- » We can implement parallel for-loops by
 - » Chunking
 - » Offset based on thread id
- » Barriers are useful to synchronize the execution

New designs/tools

- » Locks, semaphores or events can be used to signal threads / control execution
- » Thread id can also be used if something should be serial
 - » Best used in combination with the other tools

Common building blocks

- » Reusable building blocks
 - » Prefix sum/scan
 - » “Parallel for” (fork/join)
- » Can be annoying to interleave sync
- » Can be slower without interleaving
 - » Since more synchronization, which has overhead
- » Lack of barriers in Go is annoying!

Next time

- » More algorithms
 - » Searching
 - » Graphs

