

Parallel Programming

Parallel algorithms 2


Morgan Ericsson

Today

- » Graphs
- » Searching

Graphs

Remember graphs ...



... and DFS

```
1 class DFSPaths:
2     def __init__(self, g:Graph, s:int) -> None:
3         self.G = g
4         self.s = s
5         self.marked = np.zeros(self.G.V, dtype=bool)
6         self.edge_to = np.zeros(self.G.V, dtype=int)
7         self._dfs(s)
```

... and DFS

```
1 def _dfs(self:DFSPaths, v:int) -> None:
2     self.marked[v] = True
3     for w in self.G.adj(v):
4         if not self.marked[w]:
5             self._dfs(w)
6             self.edge_to[w] = v
```

Concurrent DFS?

- » We use a recursive algorithm
- » The graph is shared, but read-only
- » `marked` and `edge_to` are shared and read-write

First, iterative version!

```
1 def _dfs(self) -> None:
2     while not self.Q.empty():
3         v = self.Q.get()
4         if self.marked[v]:
5             continue
6         self.marked[v] = True
7         for w in self.G.adj(v):
8             if not self.marked[w]:
9                 self.Q.put(w)
```


First, iterative version!

```
1 class DFSPaths:
2     def __init__(self, g:Graph, s:int) -> None:
3         self.G = g
4         self.s = s
5         self.marked = np.zeros(self.G.V, dtype=bool)
6         self.Q = LifoQueue()
7         self.Q.put(self.s)
8         self._dfs()
```

Making dfs concurrent

```
1 def _dfs(self) -> None:
2     while not self.Q.empty():
3         v = self.Q.get()
4         lock.acquire()
5         if self.marked[v]:
6             continue
7         lock.release()
8         lock.acquire()
9         self.marked[v] = True
10        lock.release()
11        for w in self.G.adj(v):
12            lock.acquire()
13            if not self.marked[w]:
14                self.Q.put(w)
15            lock.release()
```

Fails

- » The queue can be empty at times, causing threads to exit prematurely
- » if ... continue inside a lock is a really bad idea
- » We can get rid of one of the if marked/if not marked checks

Second attempt

```
1  def _dfs(self) -> None:
2      while True:
3          v = self.Q.get()
4          lock.acquire()
5          mark = self.marked[v]
6          lock.release()
7          if mark:
8              continue
9          lock.acquire()
10         self.marked[v] = True
11         lock.release()
12         for w in self.G.adj(v):
13             self.Q.put(w)
```

Ugly

```
1 lock = Lock()  
2  
3 lock.acquire()  
4 try:  
5     # do something critical  
6     pass  
7 finally:  
8     lock.release()
```

We can use a context manager

```
1 lock = Lock()  
2  
3 with lock:  
4     # do something critical  
5     pass
```

Third attempt

```
1  def _dfs(self) -> None:
2      while True:
3          v = self.Q.get()
4          with lock:
5              mark = self.marked[v]
6              if not mark:
7                  self.marked[v] = True
8              if mark:
9                  continue
10         for w in self.G.adj(v):
11             self.Q.put(w)
```

Still fails

- » We have not addressed the queue
- » So it never exits
- » We can use the same trick we used for quicksort

Fourth attempt

```
1  def _dfs(self) -> None:
2      while true:
3          self.qsem.acquire()
4          v = self.Q.get()
5          with lock:
6              mark = self.marked[v]
7              if not mark:
8                  self.marked[v] = True
9              if mark:
10                 continue
11             for w in self.G.adj(v):
12                 self.Q.put(w)
13                 self.qsem.release()
```

Still fails

- » Will not exist, stuck on the queue semaphore
- » We can check if all node are processed

Fifth attempt

```
1 def _dfs(self) -> None:
2     while true:
3         self.qsem.acquire()
4         if self.cnt == self.G.V:
5             break
6         v = self.Q.get()
7         with lock:
8             mark = self.marked[v]
9             if not mark:
10                 self.marked[v] = True
11                 self.cnt += 1
12             if not mark:
13                 for w in self.G.adj(v):
14                     self.Q.put(w)
15                     self.qsem.release()
16                     if self.cnt == self.G.V:
17                         self.tsig.set()
```

Spawning threads

```
1 def __init__(self, g:Graph, nt:int) -> None:
2     self.G = g
3     self.marked = np.zeros(self.G.V, dtype=bool)
4     self.Q = LifoQueue()
5     self.lock = Lock()
6     self.tsig = Event()
7     self.cnt = 0
8
9     for v in range(self.G.V): self.Q.put(v)
10    self.qsem = Semaphore(self.G.V)
11    self.ts = [Thread(target=self._dfs) for _ in range(nt)]
12    for t in self.ts: t.start()
13    self.tsig.wait()
14    for t in self.ts: t.join()
```

One problem remains

```
1 while true:
2     self.qsem.acquire()
3     if self.cnt == self.G.V:
4         break
```

- » We previously introduced a timeout and a loop
- » How should we set the timeout?
 - » Unnecessary wake-ups or long delays
- » We can use another trick!

Spawning threads

```
1 def __init__(self, g:Graph, nt:int) -> None:
2     self.G = g
3     self.marked = np.zeros(self.G.V, dtype=bool)
4     self.Q = LifoQueue()
5     self.lock = Lock()
6     self.tsig = Event()
7     self.cnt = 0
8
9     for v in range(self.G.V): self.Q.put(v)
10    self.qsem = Semaphore(self.G.V)
11    self.ts = [Thread(target=self._dfs) for _ in range(nt)]
12    for t in self.ts: t.start()
13    self.tsig.wait()
14    self.qsem.release(nt)
15    for t in self.ts: t.join()
```

Can we improve?

```
1 with lock:
2     mark = self.marked[v]
3     if not mark:
4         self.marked[v] = True
5         self.cnt += 1
```

Few locks

- » We use a single lock for the marked array
- » It is safe to mark two different vertices in parallel
- » This can be a problem for large graphs
- » Reduced concurrency, parallelism, performance

Fix?

Setup

```
1 lock = Semaphore(10)
```

Threads

```
1 with lock:  
2     mark = self.marked[v]  
3     if not mark:  
4         self.marked[v] = True  
5         self.cnt += 1
```

No!

- » The semaphore allows 10 threads to enter the critical section
- » Which can all modify the same vertex
- » The locks must be connected to the vertices

Fix?

Setup

```
1 locks = [Lock() for _ in range(self.G.V)]
```

Threads

```
1 with locks[v]:  
2     mark = self.marked[v]  
3     if not mark:  
4         self.marked[v] = True  
5         self.cnt += 1
```

Better, not good

- » Many locks in large graphs
- » Waste of resources
- » We will most likely not access all vertices at once
- » But, an idea we can work with

Fix?

Setup

```
1 locks = [Lock() for _ in range(NUM_LOCKS)]
```

Threads

```
1 with locks[v % NUM_LOCKS]:  
2     mark = self.marked[v]  
3     if not mark:  
4         self.marked[v] = True  
5         self.cnt += 1
```

Back to Go

```
1  type DFSPaths struct {  
2      g      Graph  
3      s, nc  int  
4      marked []bool  
5      edge_to []int  
6      lock   sync.Mutex  
7      ch     chan int  
8  }
```


Back to Go

```
1 func (dfp *DFSPaths) dfs() {  
2     var mark bool  
3     for v := range dfp.ch {  
4         dfp.lock.Lock()  
5         mark = dfp.marked[v]  
6         dfp.lock.Unlock()  
7  
8         if mark { continue }  
9  
10        dfp.lock.Lock()  
11        dfp.marked[v] = true  
12        dfp.nc += 1  
13        dfp.lock.Unlock()  
14  
15        // ...
```

Back to Go

```
1      for _, w := range dfp.g.Adj(v) {
2          dfp.lock.Lock()
3          mark = dfp.marked[w]
4          dfp.lock.Unlock()
5          if !mark {
6              dfp.ch <- w
7              dfp.lock.Lock()
8              dfp.edge_to[w] = v
9              dfp.lock.Unlock()
10         }
11     }
12     dfp.lock.Lock()
13     if dfp.nc == dfp.g.V() {
14         close(dfp.ch)
15     }
16     dfp.lock.Unlock()
17 }
18 }
```


Prim's algorithm (MST)



3-5, 1

1-3, 2

4-5, 3

2-4, 4

1-2, 5

0-3, 6

0-1, 7

4-7, 8

0-6, 9

1-4, 10

3-6, 12

2-7, 13

5-6, 14

5-7, 16

6-7, 20

Prim's algorithm (MST)

- » Prim's algorithm is sequential
 - » We add one edge per iteration
 - » Which edges are available to add depends on the previous iteration
- » No need for exhaustive search, we know which edge to add
- » But, we can improve how we find that edge

Remember parallel reductions

- » We can replace the heap with a parallel reduction
- » Find the minimum weight across a large list of edges

All-pairs shortest path (APSP)

- » Dijkstra and Bellman-Ford computed single source shortest path
- » APSP is equivalent to running, e.g., Dijkstra on all possible sources
 - » Which can be done in parallel
 - » But unnecessary overhead
- » Floyd-Warshall

Remember **EWDiGraph**

- » Directed graph with edge weights
- » We can iterate over all edges
 - » Get source, destination, and weight

Distance matrix

```
1 def gendistmatrix(g):
2     dm = np.full((g.V, g.V), np.inf)
3     np.fill_diagonal(dm, 0)
4
5     for e in g.edges
6         dm[e.src, e.dst] = e.weight
7
8     return dm
```

Serial Floyd-Warshall

```
1 def FloydW(dm, V):  
2     for k in range(V):  
3         for i in range(V):  
4             for j in range(V):  
5                 dm[i, j] = min(dm[i, j], dm[i, k] + dm[k, j])
```

Concurrency

- » We cannot do the outer loop
- » We can do the i or j loops
 - » Which one?
- » Outmost is better, since that leaves more work per thread
- » Can mix in this case, just interested in all combinations of i and j .
 - » No dependencies between them, only to k

Benefits?

- » Floyd-Warshall is $O(N^3)$ operations
- » Finding min is $O(N)$
- » Parallel require slightly more operations, but executed in parallel
 - » Assume P processors
 - » Floyd-Warshall takes N^3 / P
 - » Finding min takes $N / P + p$
- » Speedup approaches P if $N \gg P$

Benefits?

- » We do not improve the number of operations required
 - » Often the opposite, we may have to do more!
 - » To achieve concurrency
- » We improve the number of operations we can do at once
 - » Parallel execution
- » Synchronization increases the number of operations and reduces the “at once”-part

Searching

Unsorted lists

- » Searching in unsorted lists is easy
 - » Just a reduction
- » E.g., find min, but looking for specific value instead

Sorted lists




```
1 def binsearch(l:list[int], x:int) -> int|None:
2     lo, hi = 0, len(l) - 1
3
4     while lo <= hi:
5         mid = (lo + hi) // 2
6         if l[mid] == x:
7             return mid
8         elif l[mid] < x:
9             lo = mid + 1
10        else:
11            hi = mid - 1
12
13    return None
```

N-array search

- » Assume N threads
- » Checks N well-spaced values, one per thread
- » Marks the values based on the result of comparisons

Example

Searching for 53



N-array search (serial)

```
1 def nary(A, lo, hi, key, intv):
2     mid, locate = [None] * (intv + 1), [None] * (intv + 2)
3     pos = None
4
5     locate[0], locate[intv + 1] = 'R', 'L'
6     while lo <= hi and pos is None:
7         mid[0] = lo - 1
8         step = (hi - lo + 1) // (intv + 1)
9
10        mark_loc()
11
12        for i in range(1, intv + 1):
13            if locate[i] != locate[i - 1]:
14                lo = mid[i - 1] + 1
15                hi = mid[i] - 1
16
17        if locate[intv] != locate[intv + 1]:
18            lo = mid[intv] + 1
19
20    return pos
```


N-array search (serial)

```
1 def mark_loc():
2     for i in range(1, intv + 1):
3         offs = step * i + (i - 1)
4         mid[i], lmid = lo + offs, lo + offs
5
6         if lmid <= hi:
7             if A[lmid] > key:
8                 locate[i] = 'L'
9             elif A[lmid] < key:
10                locate[i] = 'R'
11            else:
12                locate[i] = 'E'
13                pos = lmid
14        else:
15            mid[i] = hi + 1
16            locate[i] = 'L'
```

Concurrent version (sketch)

- » The while loop will be the starting point of the threads
- » The variables before should be done synchronized by a single thread
- » `mark_loc` will be the main concurrent work
- » The for-loop and if-statement needs to be run by a single thread

Concurrent version (sketch)

1. Spawn threads that start the while
2. Barrier
3. Do updates
4. Barrier
5. Run mark-lock (consider shared/local data)
6. barrier
7. for and if
8. barrier

