

Parallel Programming

I/O and concurrency

Morgan Ericsson

Today

- » Concurrency and I/O
- » HTTP requests
- » HTTP servers
- » REST APIs

Concurrency and I/O?

- » So far, focus on parallelism
- » Remember I/O vs CPU bound
- » Another reason to use concurrency is to “hide” I/O wait
 - » We simply do other things while we wait for I/O to finish

Concurrency and I/O?

Device	CPU cycles	Proportional
L1 cache	3	3 sec
L2 cache	14	14 sec
RAM	250	250 sec
disk	41 000 000	1.3 years
network	240 000 000	7.6 years

An example

Example settings

- » A typical example of I/O bound tasks is networking and web APIs
- » We will use the Nobel prize REST API
- » Information about prizes and laurates

Example API call

```
1 import requests
2 import json
3
4 BU = 'http://api.nobelprize.org/2.1/'
5 PU = 'nobelPrize/'
6 LU = 'laureate/'
7
8 r = requests.get(f'{BU}{PU}phy/2022')
9 if r.ok:
10     for np in r.json():
11         for ll in np['laureates']:
12             print(ll['fullName']['en'])
```

HTTP Request

- » HTTP is the hypertext transfer protocol
- » Simple protocol to transfer text
 - » GET, PUT, POST, ...
- » Used by web servers
- » And to develop API
 - » More about this in a later course

HTTP Request

```
GET /clock HTTP/1.1  
Host: localhost:3000  
User-Agent: curl/7.81.0  
Accept: */*
```

HTTP Response

HTTP/1.1 200 OK

Date: Mon, 26 Feb 2024 12:40:49 GMT

Content-Length: 29

Content-Type: text/plain; charset=utf-8

12:40:49 UTC

HTTP Request

```
POST /area HTTP/1.1
Host: localhost:3000
User-Agent: curl/7.81.0
Accept: */*
Content-Type: application/json
Content-Length: 24

{"height":30,"width":20}
```

HTTP Response

HTTP/1.1 200 OK

Date: Mon, 26 Feb 2024 12:53:20 GMT

Content-Length: 12

Content-Type: text/plain; charset=utf-8

```
{"area":600}
```

JSON

```
1  [  
2    {"id": "940",  
3     "knownName": {  
4       "en": "Michael W. Young",  
5       "se": "Michael W. Young"  
6     },  
7     "gender": "male",  
8     "birth": {  
9       "date": "1949-03-28",  
10      "place": {  
11        "city": {"en": "Miami, FL",  
12        "no": "Miami, FL",  
13        "se": "Miami, FL"  
14      },  
15      "country": {  
16        "en": "USA",  
17        "no": "USA",  
18        "se": "USA"  
19      }  
20      ...  
21    ]
```

JSON

- » JavaScript Object Notation
- » Text-based interchange format
- » Supports some types: number, string, array, boolean, object
- » Language-independent, most languages has support for json
 - » e.g. `json` in Python or `encoding/json` in Go

Any (unknown) type

```
1 var a any
2 a = "Hello"
3 b := string(a)
4 fmt.Println(b)
```

./prog.go:12:14: cannot convert a (variable of type any) to type string: need type assertion

Type assertions in Go

```
1 var a any
2 a = "Hello"
3 b := a.(string)
4 fmt.Println(b)
```


Type assertions in Go

```
1 var a any
2 a = "Hello"
3 b := a.(int)
4 fmt.Println(b)
```

panic: interface conversion: interface {} is string, not int

Example API call

```
1 url := "https://api.nobelprize.org/2.1/nobelPrize/phy/2022"
2
3 resp, err := http.Get(url)
4 if err != nil {
5     log.Fatal(err)
6 }
7 defer resp.Body.Close()
8
9 if resp.StatusCode != http.StatusOK {
10     log.Fatal(resp.StatusCode)
11 }
12
13 data, err := io.ReadAll(resp.Body)
14 if err != nil {
15     log.Fatal(err)
16 }
```

Example API call

```
1 var result []map[string]any
2 json.Unmarshal(data, &result)
3
4 tmp := result[0]["laureates"].([]any)
5 for _, k := range tmp {
6     fn := k.(map[string]any)["fullName"]
7     en := fn.(map[string]any)["en"]
8     fmt.Println(en.(string))
9 }
```

Why concurrency?

- » Assume we want to find the birth months of all laureates between 2000 and 2022
- » We use something similar to the previous example
- » One call per field and year, and then another per laureate
 - » about 400 calls the API
- » Can take minutes, most of it waiting for I/O

Creating a server

A server implementation in Go

```
1 func getTime(w http.ResponseWriter, r *http.Request) {  
2     ct := time.Now()  
3     io.WriteString(w, ct.Format(time.RFC1123))  
4 }  
5  
6 func main() {  
7     http.HandleFunc("/time", getTime)  
8     http.ListenAndServe(":3000", nil)  
9 }
```

And a client in Python

```
1 import requests
2
3 r = requests.get('http://localhost:3000/time')
4 if r.ok:
5     print(f'The time is {r.text}')
```

Latency

- » As discussed, there is some latency in calling the API
- » Same machine, about 0.002 seconds per call
- » Same network, about 0.014 seconds per call

Latency

- » We can create a server that simulates latency
- » Latency from, e.g.,
 - » computation
 - » networks
 - » databases and other server I/O
 - » server load
 - » ...

Latency

```
1 func simLatency(w http.ResponseWriter, r *http.Request) {  
2     time.Sleep(2*time.Second)  
3     io.WriteString(w, "ok")  
4 }  
5  
6 // ...  
7  
8 http.HandleFunc("/latency", simLatency)
```

No surprise

- » Unsurprisingly, 100 calls to the API take about 200 seconds
- » If we use ten threads to make the calls, it takes about 20 seconds
 - » Surprising?

Threaded client

```
1 def make_req(n):
2     for i in range(n):
3         r = requests.get('http://localhost:3000/latency')
4         if not r.ok:
5             print("req failed...", r.status_code)
6
7 ts = [Thread(target=make_req, args=(10,)) for i in range(10)]
8 for t in ts:
9     t.start()
10 for t in ts:
11     t.join()
```

What did we learn?

- » Remember the Python GIL?
 - » no problem, since threads are mostly blocked waiting for I/O
- » Our server can handle ten concurrent requests
 - » the handler functions run as goroutines...
 - » and we mainly sleep while serving the request
 - » which is I/O wait

Can we do 100?

- » Yes, 100 threads take a little more than two seconds
- » If we fake work instead of sleeping? $100\,000^2$ multiplications per call
- » About 300 seconds for serial, 30 for ten threads, and 6 for 100 threads
 - » Remember, server compute is not infinite
 - » So, do not hammer servers (more later)

Creating an API

A simple API to calculate the area

- » Pass height and width to an endpoint
- » Compute the area and return the value
- » Encoded as JSON

Structs and tags

```
1  type Dimensions struct {  
2      Height int `json:"height"`  
3      Width  int `json:"width"`  
4  }  
5  
6  type Area struct {  
7      Area int `json:"area"`  
8  }
```

Checking for the method

```
1 func CalcArea(w http.ResponseWriter, r *http.Request) {  
2     switch r.Method {  
3     case "POST":  
4         w.WriteHeader(http.StatusOK)  
5     case "GET":  
6         w.WriteHeader(http.StatusOK)  
7     default:  
8         w.WriteHeader(http.StatusNotFound)  
9     }  
10 }
```

Handling a post

- » Find JSON in the body
- » “Decode” to a struct
- » Compute the area
- » “Encode” the area struct
- » Write the response

Handling a post

```
1 case "POST":
2     var dims Dimensions
3     data, _ := io.ReadAll(r.Body)
4
5     if err := json.Unmarshal(data, &dims); err == nil {
6         area := Area{dims.Height * dims.Width}
7         if res, err := json.Marshal(area); err == nil {
8             w.Write(res)
9         }
10    }
```

BAD!

- » Works if everything is “perfect”
- » Cannot assume that!
- » So, not enough error handling!

Client for post

```
1 import requests
2 import json
3
4 vals = {'height':10, 'width':20}
5 r = requests.post('http://localhost:3000/area', json=vals)
6 if r.ok:
7     jo = r.json()
8     print(f'The area is {jo["area"]}')

```

Handling a get

- » Extract the query params from the URL
 - » . . . ?a=b&c=d& . . .
- » ... (same as post)

Handling a get

```
1 case "GET":
2     height, _ := strconv.Atoi(r.URL.Query().Get("height"))
3     width, _ := strconv.Atoi(r.URL.Query().Get("width"))
4     area := Area{height * width}
5
6     if res, err := json.Marshal(area); err == nil {
7         w.Write(res)
8     }
```


Client for get

```
1 import requests
2 import json
3
4 vals = {'height':10, 'width':20}
5 url = f'http://localhost:3000/area?height={vals["height"]}&width={vals["width"]}'
6 r = requests.get(url)
7 if r.ok:
8     jo = r.json()
9     print(f'The area is {jo["area"]}')

```

Client for get

```
1 import requests
2 import json
3
4 vals = {'height':10, 'width':20}
5 r = requests.get('http://localhost:3000/area', data=vals)
6 if r.ok:
7     jo = r.json()
8     print(f'The area is {jo["area"]}')

```

