

Parallel Programming

I/O and concurrency 2

Morgan Ericsson

Today

- » Context
- » Useful patterns
- » HTTP servers
- » REST APIs

Remember the API to compute area

```
1 func CalcArea(w http.ResponseWriter, r *http.Request) {  
2     switch r.Method {  
3     case "POST":  
4         w.WriteHeader(http.StatusOK)  
5     case "GET":  
6         w.WriteHeader(http.StatusOK)  
7     default:  
8         w.WriteHeader(http.StatusNotFound)  
9     }  
10 }
```

Remember the API to compute area

```
1 case "POST":
2     var dims Dimensions
3     data, _ := io.ReadAll(r.Body)
4
5     if err := json.Unmarshal(data, &dims); err == nil {
6         area := Area{dims.Height * dims.Width}
7         if res, err := json.Marshal(area); err == nil {
8             w.Write(res)
9         }
10    }
```

And the Python client

```
1 import requests
2 import json
3
4 vals = {'height':10, 'width':20}
5 url = f'http://localhost:3000/area?height={vals["height"]}&width={vals["width"]}'
6 r = requests.get(url)
7 if r.ok:
8     jo = r.json()
9     print(f'The area is {jo["area"]}')
```

Wrap in function

- » Assume we want to port the client to Go
- » And wrap it inside a function, e.g.,
 - » `func CalcArea(w, h int) (int, error)`
`{}`
- » Exposes many issues/new ideas!

Simple client

```
1 url := "http://localhost:3000/area"
2 body := []byte(`{"height":30, "width":20}`)
3
4 resp, err := http.Post(url, "application/json", bytes.NewBuffer(body))
5 if err != nil {
6     log.Fatal("Request failed:", err)
7 }
8 defer resp.Body.Close()
9
10 if resp.StatusCode != http.StatusOK {
11     log.Fatal(resp.Status)
12 }
```

Simple client

```
1 data, err := io.ReadAll(resp.Body)
2 if err != nil {
3     log.Fatal("Error reading response")
4 }
5
6 // We should use the Area type and unmarshal to that, but...
7 var result map[string]int
8 json.Unmarshal(data, &result)
9 fmt.Println("The area is", result["area"])
```


Wrapping in a function

- » How should we handle errors?
 - » retries, timeouts, ...
- » How should we act responsibly?
 - » and not kill the server
- » Blocking or non-blocking?

Context

- » A context allows for “per request” data and functions
- » An easy way to cancel requests
 - » time-out / deadline
- » Parent and child contexts
 - » Values in tree
 - » Cancels downwards

Context

```
1 func FunWithContext(ctx context.Context) {  
2     // Do something with the context  
3 }
```

Values

```
1 func FunWithContext(ctx context.Context) {
2     if val := ctx.Value("mydata"); val != nil {
3         // The value is of type any, so
4         // we need a type assertion
5         fmt.Println("Value:", val.(string))
6     }
7 }
8
9 func main() {
10     c := context.Background()
11     FunWithContext(c)
12 }
```

Values

```
1 func FunWithContext(ctx context.Context) {
2     if val := ctx.Value("mydata"); val != nil {
3         fmt.Println("Value:", val.(string))
4     }
5 }
6
7 func main() {
8     c := context.WithValue(context.Background(), "mydata", "abc123")
9
10    FunWithContext(c)
11 }
```

Values

```
1 func FunWithContext(ctx context.Context) {
2     if val := ctx.Value("more"); val != nil {
3         fmt.Println("Value:", val.(int))
4     }
5 }
6
7 func main() {
8     c := context.WithValue(context.Background(), "mydata", "abc123")
9     c = context.WithValue(c, "more", 123)
10    FunWithContext(cc)
11 }
```

Remember the done channel

```
1 func doSmtH(done <-chan interface{}) {  
2     for {  
3         select {  
4             case <-done:  
5                 return  
6             default:  
7                 }  
8  
9         // Do some work  
10    }  
11 }
```

Remember the done channel

- » Same functionality in the context
 - » either explicit or timed cancel

Same but context

```
1 func doSmtH(ctx context.Context) {  
2     for {  
3         select {  
4             case <-ctx.Done():  
5                 return  
6             default:  
7                 }  
8  
9         // Do some work  
10    }  
11 }
```

Same but context

```
1 func main() {  
2     c, cancel := context.WithCancel(context.Background())  
3     go doSmt(c)  
4     cancel()  
5 }
```

Timeout

```
1 func main() {
2     var wg sync.WaitGroup
3     c, cancel := context.WithTimeout(context.Background(), 2*time.Second)
4
5     wg.Add(1)
6     go func() {
7         defer wg.Done()
8         defer cancel() // Should cancel if no timeout to cleanup
9         doSmth(c)
10    }()
11
12    wg.Wait()
13    fmt.Println("Done!")
14 }
```

Deadline

- » Similar to Timeout, but absolute time rather than offset
 - » `WithTimeout(..., 2 * time.Second)`
 - » `WithDeadline(..., time.Now().Add(2 * time.Second))`

Parents and children

- » All the childrens' done channels are closed when the parent's done channel is closed
 - » so, parent with shorter timeout/deadline will cancel children before their timeout/deadline
- » Cancelling a child does not cancel the parent

Example

```
1 func main() {  
2     c1, cf1 := context.WithCancel(context.Background())  
3     c2, cf2 := context.WithCancel(c1)  
4     c3, cf3 := context.WithCancel(c2)  
5  
6     // cf3() cancels things called with c3  
7     // cf2() cancels things called with c3 and c2  
8     // cf1() cancels things called with c3, c2, and c1  
9 }
```

Example

```
1 func main() {
2     var wg sync.WaitGroup
3     c, _ := context.WithTimeout(context.Background(), 2*time.Second)
4     cc, cancel2 := context.WithTimeout(c, 5*time.Second)
5
6     wg.Add(1)
7     go func() {
8         defer wg.Done()
9         defer cancel2()
10        doSmtH(cc)
11    }()
12
13    wg.Wait()
14    fmt.Println("Done!")
15 }
```

Remember SimLatency

```
1 func simLatency(w http.ResponseWriter, r *http.Request) {
2     time.Sleep(15*time.Second)
3     io.WriteString(w, "ok")
4 }
5
6 func main() {
7     // ...
8
9     http.HandleFunc("/latency", simLatency)
10    http.ListenAndServe(":3000", nil)
11 }
```


Can we get the client to time out?

- » Yes, contexts!
- » But we need to pass a context with the request
- » `NewRequestWithContext()`
- » We will do it in two steps,
 - » first, change to request and client
 - » then, add timeout

Previous version

```
1 url := "http://localhost:3000/simlatency"
2
3 resp, err := http.Get(url)
4 if err != nil {
5     log.Fatal("Request failed:", err)
6 }
7 defer resp.Body.Close()
8
9 if resp.StatusCode != http.StatusOK {
10     log.Fatal(resp.Status)
11 }
```

Request and client

```
1 url := "http://localhost:3000/latency"
2
3 req, err := http.NewRequest("GET", url, nil)
4 if err != nil {
5     log.Fatal("Error creating the request")
6 }
7
8 client := &http.Client{}
9 resp, err := client.Do(req)
10 if err != nil {
11     log.Fatal("Request failed")
12 }
13
14 // ...
```

With timeout

```
1 url := "http://localhost:3000/latency"
2 c, cancel := context.WithTimeout(context.Background(), 5*time.Second)
3
4 req, err := http.NewRequestWithContext(c, "GET", url, nil)
5 if err != nil {
6     if errors.Is(err, context.DeadlineExceeded) {
7         log.Fatal("Timeout")
8     } else {
9         log.Fatal("Error creating the request")
10    }
11 }
12
13 client := &http.Client{}
14 resp, err := client.Do(req)
15 if err != nil {
16     log.Fatal("Request failed")
17 }
18
19 // ...
```

We can of course cancel...

```
1 // ...
2 go func() {
3     defer cancel()
4     time.Sleep(2 * time.Second)
5 }()
6
7 client := &http.Client{}
8 resp, err := client.Do(req)
9 if err != nil {
10     switch {
11     case errors.Is(err, context.DeadlineExceeded):
12         log.Fatal("Timeout")
13     case errors.Is(err, context.Canceled):
14         log.Fatal("Canceled")
15     default:
16         log.Fatal(err)
17     }
18 }
```

Back to the **CalcArea** function

```
1 func CalcArea(w, h int) (int, error) {  
2     // do the request  
3     // extract and return the area  
4     // or a sane error  
5 }  
6  
7 func main() {  
8     if area, ok := CalcArea(20, 30); ok == nil {  
9         fmt.Printf("The area is %d", area)  
10    }  
11 }
```

Wait...

- » Did we actually fix anything?
- » Remember,
 - » I/O is slow
 - » blocking can be bad
- » New CalcArea blocks waiting for I/O
- » We know how to fix that!

Adding channels

```
1 func main() {
2     ch := make(chan int)
3     go func() {
4         if area, ok := CalcArea(20, 30); ok == nil {
5             ch <- area
6         }
7         close(ch)
8     }()
9
10    if a, ok := <-ch; ok {
11        fmt.Printf("The area is %d", a)
12    }
13 }
```


Adding channels

- » We love channels
- » But they can be annoying
- » And a bit verbose

Async python

```
1  async def main():
2      pl = {'height':20, 'width':30}
3      url = 'http://localhost:3000/area'
4      async with ClientSession() as session:
5          async with session.post(url, json=pl) as resp:
6              # Do stuff!
7              res = await resp.json()
8              print("The area is", res["area"])
9
10  asyncio.run(main())
```

Emulate in Go?

```
1 type Future[T any] interface {  
2     Result() (T, error)  
3 }
```

- » We define a type for futures
- » Contains a single method that blocks waiting for the result (and error)
- » Similar to futures in Python or Java

InnerFuture

```
1  type InnerFuture[T any] struct {  
2      once sync.Once  
3      wg sync.WaitGroup  
4  
5      res T  
6      err error  
7      resCh <-chan T  
8      errCh <-chan error  
9  }
```

Result

```
1 func (f *InnerFuture[T]) Result() (T, error) {
2     f.once.Do(func() {
3         f.wg.Add(1)
4         defer f.wg.Done()
5         f.res = <-f.resCh
6         f.err = <-f.errCh
7     })
8
9     f.wg.Wait()
10
11     return f.res, f.err
12 }
```

- » The blocking result is implemented using channels
- » Can be called multiple times, only waits the first time (`once.Do`)

New CalcArea

```
1 func CalcArea(h, w int) Future[int] {
2     resCh := make(chan int)
3     errCh := make(chan error)
4
5     go func() {
6         // Do all the stuff
7
8         resCh <- // res here (if any)
9         errCh <- // error here (if any)
10    }()
11
12    return &InnerFuture[int]{resCh: resCh, errCh: errCh}
13 }
```

- » The new `CalcArea` returns an `InnerFuture`
- » That we can “wait for” by calling `result`

Calling

```
1 future := CalcArea(20, 30)
2 // do stuff
3
4 if res, err := future.Result(); res == nil {
5     fmt.Printf("The area is %d", res)
6 }
```

Retries

- » What should we do if there is a failure?
- » Try multiple times? In a loop?
- » Probably not a great idea
 - » Some errors will probably not fix themselves
 - » And if they will, maybe not in fractions of seconds

Retries

```
1 type WorkFn func(context.Context) (string, error)
2
3 func Retry(fn WorkFn, retries int, delay time.Duration)
```

Retries

```
1 func Retry(fn WorkFn, retries int, delay time.Duration) WorkFn {
2     return func(ctx context.Context) (string, error) {
3         for r := 0; ; r++ {
4             resp, err := fn(ctx)
5             if err == nil || r >= retries {
6                 return resp, err
7             }
8
9             select {
10            case <-time.After(delay):
11            case <-ctx.Done():
12                return "", ctx.Err()
13            }
14        }
15    }
16 }
```

Retries

```
1 r := Retry(MyWorker, 5, 2*time.Second)
2 res, err := r(context.Background())
```

Throttle

- » We should not hammer the server
- » There is often an acceptable rate for calls
 - » Calls per time window
- » Can figured out via headers and error 429 (too many requests)

Throttle

- » Can be implemented in different ways
 - » E.g., sleep between calls
- » We will use token bucket

Token bucket

- » Bucket contains tokens
- » To make a call, we take a token from the bucket
- » The bucket has a capacity
- » Tokens are refilled over time

The refiller

```
1 refiller := func() {
2     ticker := time.NewTicker(d)
3     go func() {
4         defer ticker.Stop()
5         for {
6             select {
7             case <-ctx.Done():
8                 return
9             case <-ticker.C:
10                t := tokens + refill
11                if t > max { t = max }
12                tokens = t
13            }
14        }
15    }
16 }
```

Throttle

```
1 func Throttle(fn WorkFn, max int, refill int, d time.Duration) WorkFn {
2     var tokens = max
3     var once sync.Once
4
5     return func(ctx context.Context) (string, error) {
6         if ctx.Err() != nil {
7             return "", ctx.Err()
8         }
9
10        once.Do(refiller)
11
12        if tokens <= 0 {
13            return "", errors.New("too many calls")
14        }
15        tokens--
16        return fn(ctx)
17    }
18 }
```


Back to the server

- » We can use many of the tricks discussed on the server as well
- » In some cases, a bit more complicated
- » For example, we need to create, populate, propagate the context

Back to the server

```
1 func simLatency(w http.ResponseWriter, r *http.Request) {  
2     time.Sleep(15*time.Second)  
3     io.WriteString(w, "ok")  
4 }  
5  
6 func main() {  
7     // ...  
8  
9     http.HandleFunc("/latency", simLatency)  
10    http.ListenAndServe(":3000", nil)  
11 }
```

Back to the server

```
1 type Handler interface {  
2     ServeHTTP(ResponseWriter, *Request)  
3 }
```

- » `http.Handler` is an interface
- » A `Server` has a handler that it can call

Back to the server

```
1 func main() {  
2     mux := http.NewServeMux()  
3  
4     mux.HandleFunc("/latency", simLatency)  
5     srv := &http.Server{Handler:mux, Addr:"0.0.0.0:3000"}  
6     srv.ListenAndServe()  
7 }
```

Context 1

- » `Server` has a field, `BaseContext`
- » The base context is a function that is used to create the context for each request
- » If nil, `context.Background()` is used

Stupid example

```
1 ctx, _ := context.WithCancel(context.Background())
2 srv := &http.Server{
3     Handler: mux,
4     Addr: "0.0.0.0:3000",
5    BaseContext: func(l net.Listener) context.Context {
6         ctx = context.WithValue(ctx, "SleepTime", 15)
7         return ctx
8     },
9 }
```

Stupid example

```
1 func simLatency(w http.ResponseWriter, r *http.Request) {  
2     ctx := r.Context()  
3     st := ctx.Value("SleepTime").(int)  
4     time.Sleep(st*time.Second)  
5     io.WriteString(w, "ok")  
6 }
```

Mux and ServeHTTP

- » Since `ServeMux` can be a handler, it must have a `ServeHTTP` function
- » This method basically finds the right handler for the route and calls `ServeHTTP` on that handler
- » The handler is specified via the `HandleFunc` function

Really?

```
1 func simLatency(w http.ResponseWriter, r *http.Request) {  
2     time.Sleep(15*time.Second)  
3     io.WriteString(w, "ok")  
4 }
```

» Where is the `ServeHTTP` function?

Really

```
1 type HandlerFunc func(ResponseWriter, *Request)
2
3 func (f HandlerFunc) ServeHTTP(w ResponseWriter, r *Request) {
4     f(w, r)
5 }
```

Chaining handler functions

```
1 func mw1(next http.Handler) http.Handler {  
2     return http.HandlerFunc(func(w http.ResponseWriter, r *http.Request) {  
3         ctx := r.Context()  
4         ctx = context.WithValue(ctx, "turtles", true)  
5         next.ServeHTTP(w, r.WithContext(ctx))  
6     })  
7 }
```

Chaining handler functions

```
1 mux.Handle("/latency", mw1(http.HandlerFunc(simLatency)))
```

Chaining handler functions

```
1 func simLatency(w http.ResponseWriter, r *http.Request) {  
2     ctx := r.Context()  
3     turtles := ctx.Value("turtles").(bool)  
4     if turtles {  
5         fmt.Println("All the way down...")  
6     }  
7     time.Sleep(15*time.Second)  
8     io.WriteString(w, "ok")  
9 }
```

So, what do we do with the context?

- » Keep information that is specific to the request
- » Should not “live” before or after the request
- » Usually used to “decorate” the request
 - » Checking authentication
 - » Extracting useful values
 - » ...

An example

A key-value store

- » Assume we want to create a server that allows us to store and retrieve values
- » We use a REST-style interface
- » A `/kvs/<value>` end-point, where `<value>` is the key we want to do something with
 - » PUT to set/update
 - » GET to retrieve
 - » DELETE to remove

Helper functions

- » We assume that there are methods Put, Get, and Delete, that operate on the underlying data structure
- » This allows us to modify how the data is stored without changing the handlers

Helper functions

```
1  var kvs = make(map[string]string)
2
3  var ErrorNoSuchKey = errors.New("No such key")
4
5  func Get(key string) (string, error) {
6      if value, ok := kvs[key]; ok {
7          return value, nil
8      } else {
9          return "", ErrorNoSuchKey
10     }
11 }
```

Handlers

- » We use the endpoint `/kvs/` followed by the name of the key
- » How can we extract the key name from the URL in the request?
- » Go 1.22 adds new ways to define endpoints that help
 - » `{key}` represents a variable that can be accessed via the request
 - » The method can be defined together with the endpoint

Handlers

```
1 mux := http.NewServeMux()  
2  
3 mux.HandleFunc("GET /kvs/{key}", KVSGet)
```

- » We define a handler for all GET requests to `/kvs/` followed by a key name
 - » e.g., `GET /kvs/my-key`
- » Note that “key” is just a name

Implementing the GET handler

```
1 func KVSGet(w http.ResponseWriter, r *http.Request) {
2     key := r.PathValue("key") // Getting the value from the request path
3
4     value, err := Get(key) // We use the helper to get a value or an error
5     if err != nil {
6         if errors.Is(err, ErrorNoSuchKey) {
7             http.Error(w, err.Error(), http.StatusNotFound)
8         } else {
9             http.Error(w, err.Error(), http.StatusInternalServerError)
10        }
11        return
12    }
13
14    w.Write([]byte(value)) // If no error, send the value
15 }
```

And a PUT (to make testing easier)

```
1 func KVSPut(w http.ResponseWriter, r *http.Request) {
2     key := r.PathValue("key")
3
4     value, err := io.ReadAll(r.Body) // Get the value for the key
5     defer r.Body.Close()
6
7     if err != nil { // No value in the request
8         http.Error(w, err.Error(), http.StatusInternalServerError)
9         return
10    }
11
12    if err = Put(key, string(value)); err != nil { // Setting key failed
13        http.Error(w, err.Error(), http.StatusInternalServerError)
14        return
15    }
16
17    w.WriteHeader(http.StatusCreated) // Key was set
18 }
```

Aside: Testing

- » We can implement a Go or Python client to test
- » We can also use various tools, e.g., curl or wget
- » We can set the key *aa* to “Hello!”
 - » `curl -X PUT -d 'Hello!' -v http://localhost:3000/kvs/aa`
- » And fetch it
 - » `curl -v http://localhost:3000/kvs/aa`

Delete for completeness

```
1 func KVSDel(w http.ResponseWriter, r *http.Request) {
2     key := r.PathValue("key")
3
4     err := Delete(key)
5     if err != nil {
6         http.Error(w, err.Error(), http.StatusInternalServerError)
7         return
8     }
9
10    w.WriteHeader(http.StatusOK)
11 }
```


The main function

```
1    mux := http.NewServeMux()  
2  
3    mux.HandleFunc("PUT /kvs/{key}", KVSPut)  
4    mux.HandleFunc("GET /kvs/{key}", KVSGet)  
5    mux.HandleFunc("DELETE /kvs/{key}", KVSDel)  
6  
7    http.ListenAndServe(":3000", mux)
```

Persistent storage?

- » The keys are currently stored in memory only
- » If the server terminates, all the keys are gone
- » Multiple ways we can fix that
 - » we will add a database for persistent storage
 - » mainly to show how to use databases in Go

Databases in Go

- » Drivers for various databases, e.g., MySQL
 - » “github.com/go-sql-driver/mysql”
- » Distributed as Go packages
- » Use “database/sql” for API

Connecting to a database

```
1 var db *sql.DB
2 db, err := sql.Open("mysql", "kvsadm@/kvs")
3 if err != nil {
4     log.Fatal(err)
5 }
```

Testing the connection

» We can use Ping to test the connection `err := db.Ping()`

Rewriting the helpers

- » We implemented the main functionality using helper functions
- » We can easily modify these to use a database instead
 - » Global variable for easy access
 - » Initialize in the main method

Rewriting the helpers

```
1 func Get(key string) (string, error) {  
2     row := db.QueryRow("SELECT value FROM kvs WHERE key = ?", key)  
3  
4     var res string  
5     if err := row.Scan(&res); err == nil {  
6         return res, ""  
7     } else {  
8         return "", ErrorNoSuchKey  
9     }  
10 }
```

Query

- » QueryRow returns zero or one rows
 - » we use Scan to find the value in the row
 - » this is where any errors (e.g., no rows) are triggered
- » If we expect multiple rows, we should use Query
- » Which returns rows and a possible error
- » Use Next on the rows to iterate over the result
- » And Scan to get the results

Contexts

- » We can pass a context to most of the database functions
 - » E.g., QueryContext, QueryRowContext and PingContext
- » Can be used to, e.g., cancel the operation

Rewriting the helpers

```
1 func Delete(key string) error {  
2     res, err := db.Exec("DELETE FROM vks WHERE key = ?", key)  
3     return err  
4 }
```

- » We use Exec since we do not expect any result from the delete query
- » We could check the returned result for, e.g., rows affected

